

CS2040S

Data Structures and Algorithms

Shortest Paths!

Midterm Scripts Published

- You should have received your script by e-mail.
- Check Luminus Gradebook! These are the official score.
- If there is a mistake (e.g., an OCR error, or incorrect transfer to Luminus), send Prof. Ben Leong an e-mail by **Wed. 23:59**.
- After Wednesday, grades are final.

Midterm Scripts Published

- You should have received your script by e-mail.
- Check Luminus Gradebook! These are the official score.
- If there is a mistake (e.g., an OCR error, or incorrect transfer to Luminus), send Prof. Ben Leong an e-mail by **Wed. 23:59**.
- After Wednesday, grades are final.

Speed Demon Competition

Announcing the winners!



Speed Demon Competition

CS2040S Speed Demon Leaderboard

Name	Runtime
Tan Kel Zin	2.234s
Cai Kai'An	2.24s
Tee Weile Wayne	2.252s
Ryan Chung Yi Sheng	2.291s
Tan Weiu Cheng	2.358s
Zhang Shichen	2.369s
Rohit Rajesh Bhat	2.376s
Lee Xiong Jie, Isaac	2.546s
Marcus Tang Xin Kye	2.789s
Lee Zheng Han	2.792s




Last Week

Introduction to Graphs

- What is a graph? What is a directed graph?
- Modelling problems as graphs.
- Graph representations.
- Graph searching: BFS/DFS

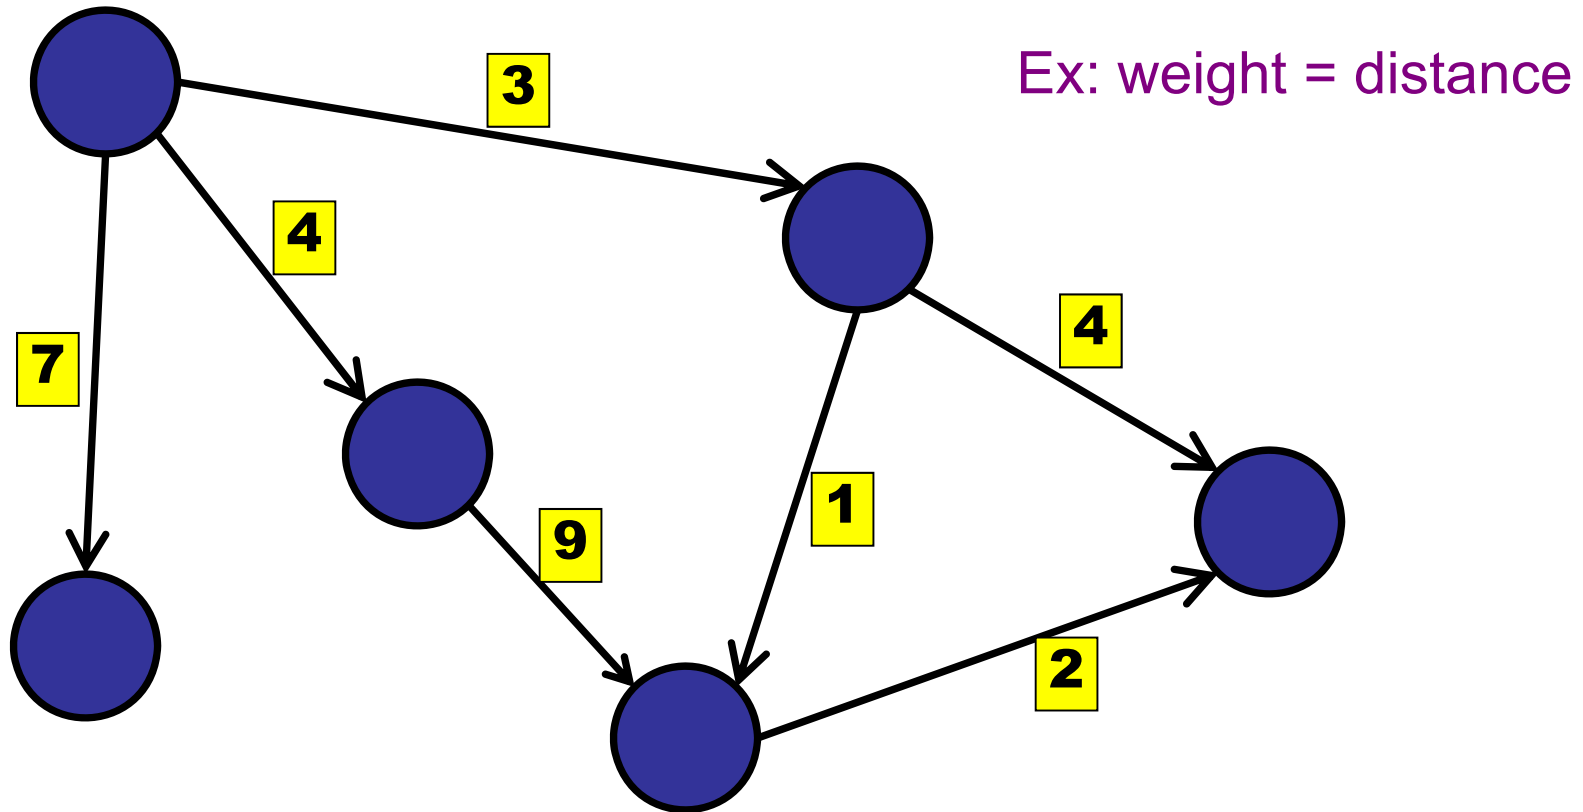
Postponed: topologic order, topological sort

SHORTEST PATHS



Weighted Graphs


Edge weights: $w(e) : E \rightarrow \mathbb{R}$



Adjacency list: stores weights with edge in NbrList


Shortest Paths

Distance from source?




What is the distance from S to D?

- 1. 2
- 2. 4
- ✓ 3. 6
- 4. 7
- 5. 9
- 6. Infinite



Shortest Paths


Distance from source?



Shortest Paths


Questions:

- How far is it from S to D?
- What is the shortest path from S to D?
- Find the shortest path from S to every node.
- Find the shortest path between every pair of nodes.



Shortest Paths


Common mistake: "Why can't I use BFS?"



Remember: BFS does not explore every path in the graph!

Shortest Paths


Common mistake: “Why can’t I use BFS?”



Remember: BFS does not explore every path in the graph!

Shortest Paths


Common mistake: "Why can't I use BFS?"



BFS finds minimum number of **HOPS** not minimum **DISTANCE**.

Shortest Paths


Notation: $\delta(u,v)$ = distance from u to v



Shortest Paths

Key idea: triangle inequality

$$\delta(S, C) \leq \delta(S, A) + \delta(A, C)$$




Shortest Paths

Maintain estimate for each distance:

```
int[] dist = new int[V.length];
```

```
Arrays.fill(dist, INFTY);
```


```
dist[start] = 0;
```



Shortest Paths

Maintain estimate for each distance:


- Reduce estimate
- Invariant: estimate \geq distance



Shortest Paths


Maintain estimate for each distance:

$\text{relax}(S, A)$



Shortest Paths


```
relax(int u, int v){  
    if (dist[v] > dist[u] + weight(u,v))  
        dist[v] = dist[u] + weight(u,v);  
}
```



Shortest Paths

Maintain estimate for each distance:


$\text{relax}(S, A)$



Shortest Paths

Maintain estimate for each distance:


$\text{relax}(A, C)$



Shortest Paths

Maintain estimate for each distance:


$\text{relax}(A, C)$



Shortest Paths

Maintain estimate for each distance:


$\text{relax}(A, B)$



Shortest Paths

Maintain estimate for each distance:


$\text{relax}(S, B)$



Shortest Paths


Maintain estimate for each distance:

$\text{relax}(B, C)$



Shortest Paths

```
for (Edge e : graph)
    relax(e)
```

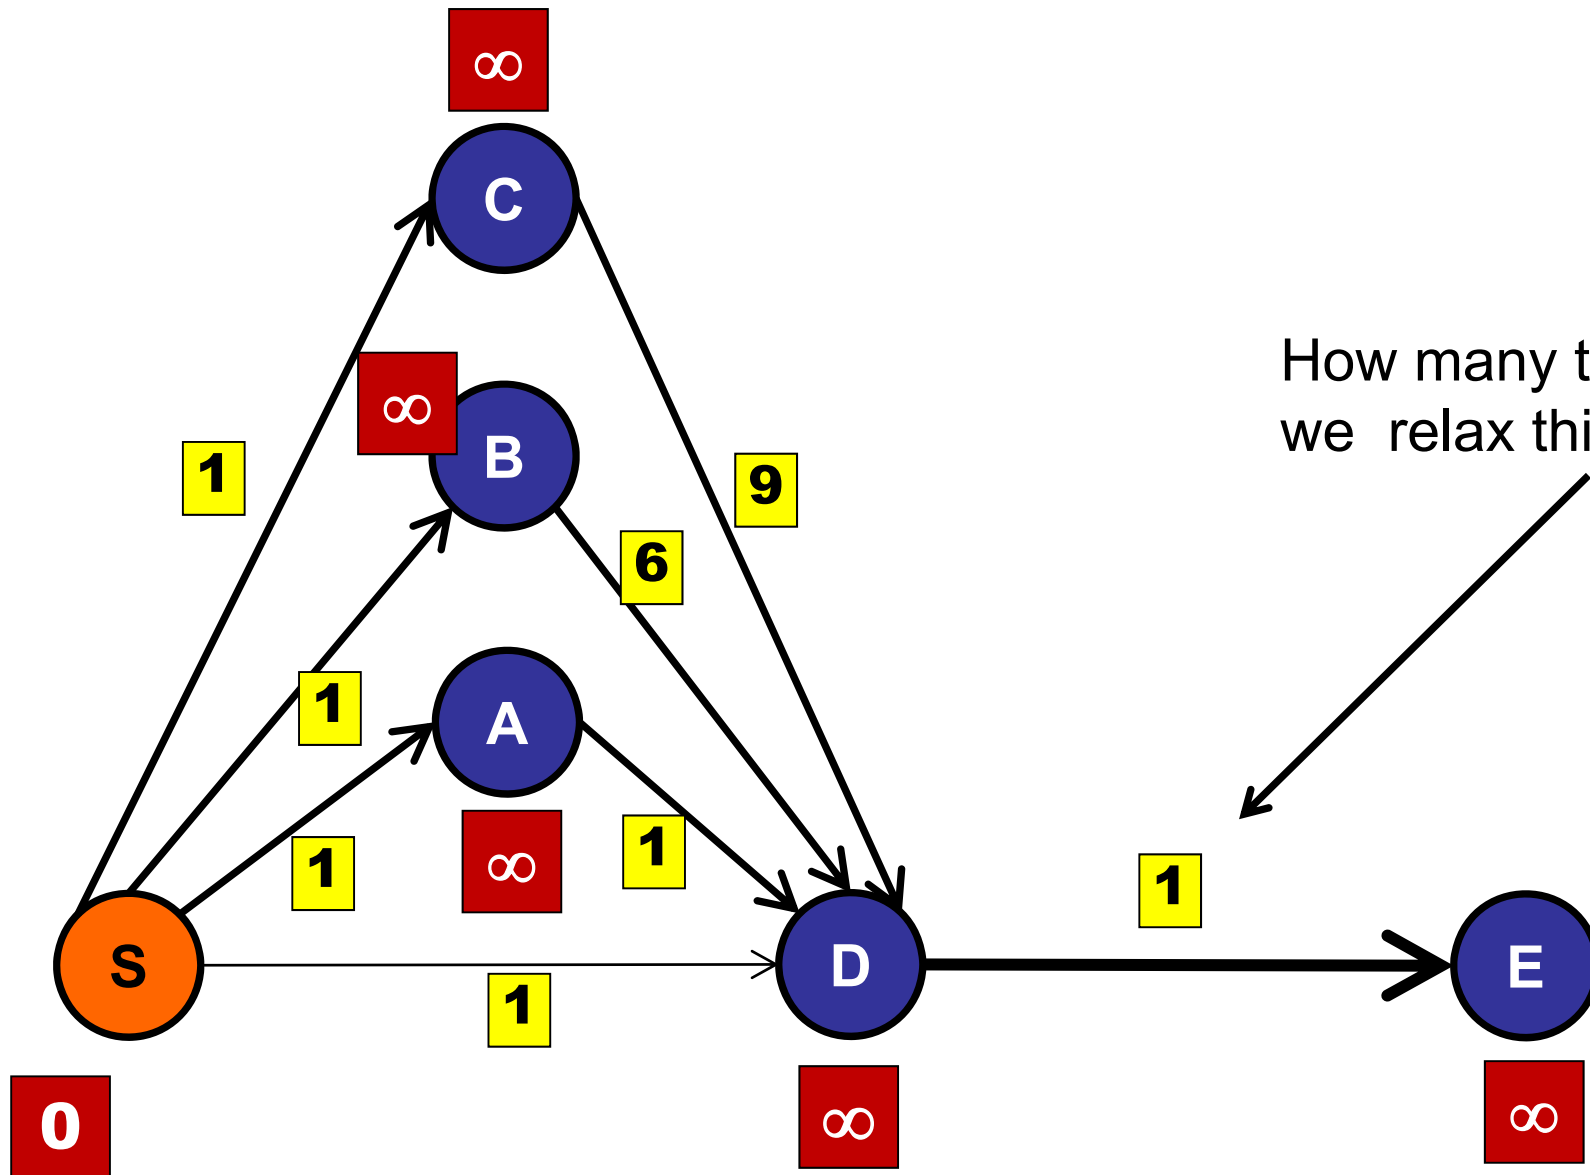


Does this algorithm work:


for every edge e : relax(e)

1. Yes
2. No
3. Only on trees.

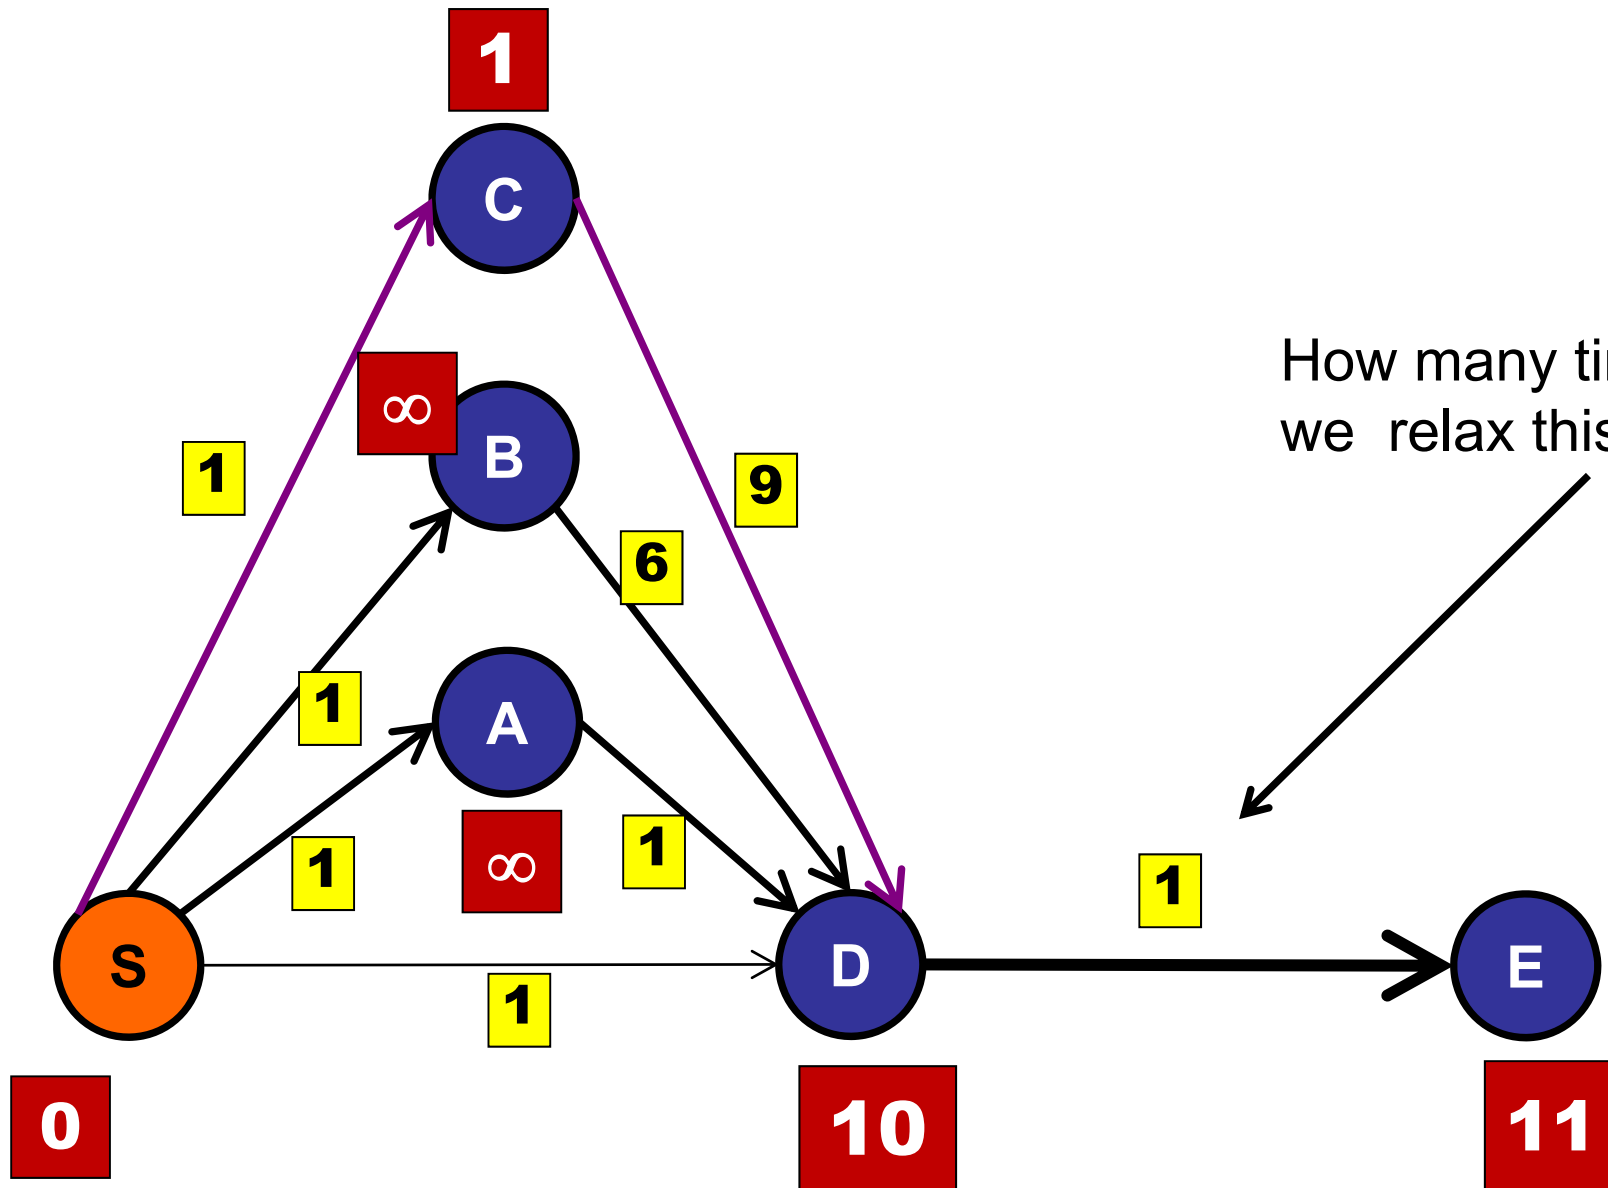
Shortest Paths



How many times might we relax this edge?

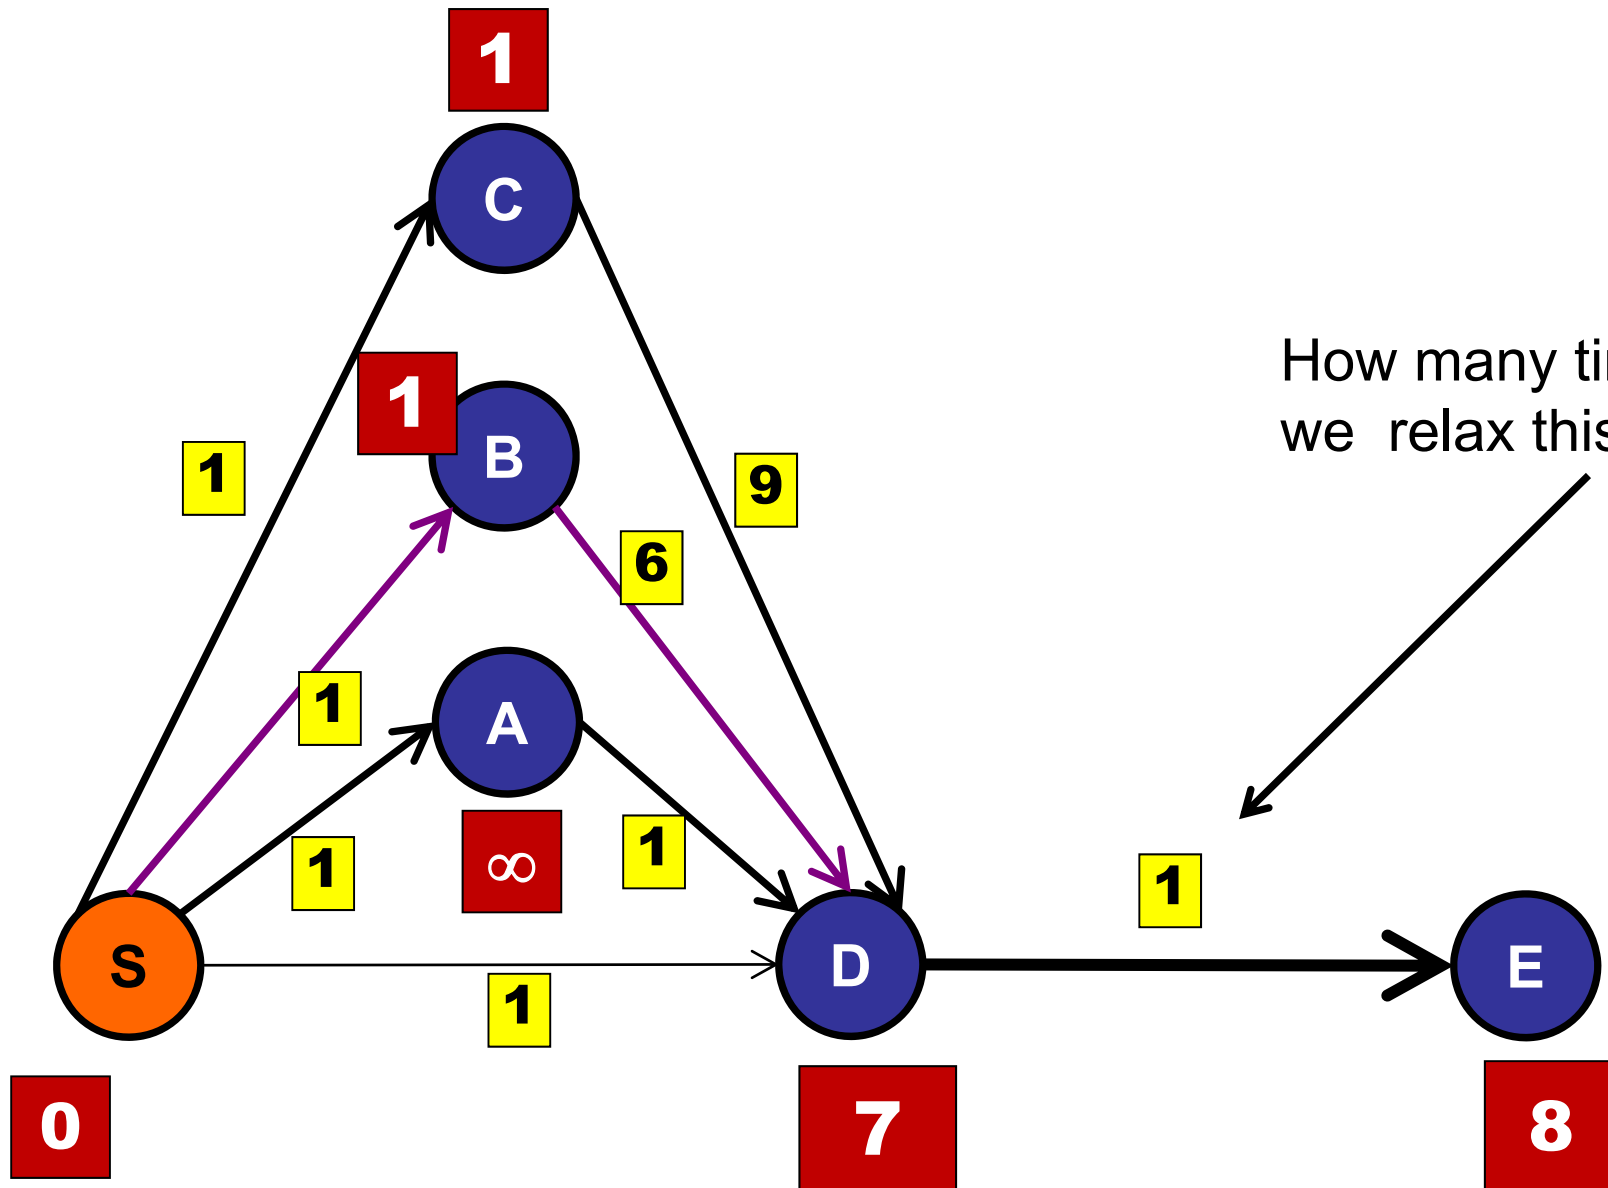


Shortest Paths



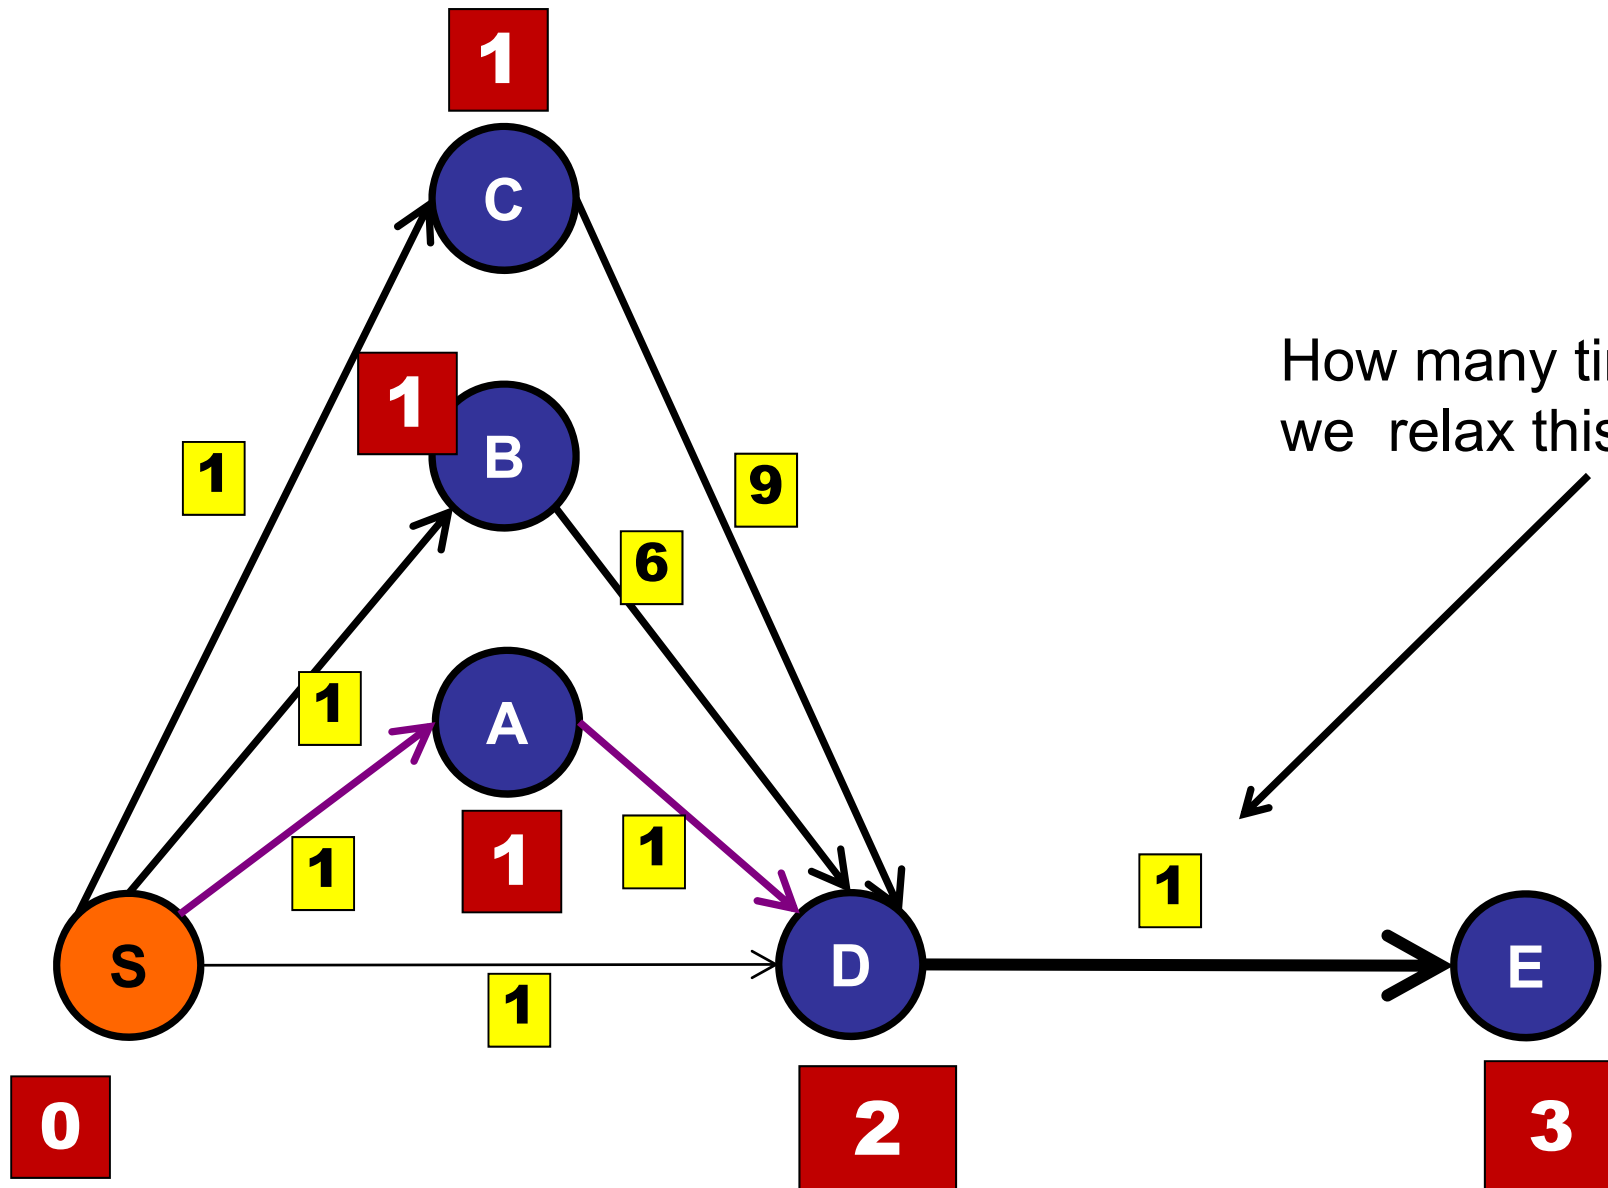
How many times might we relax this edge?

Shortest Paths

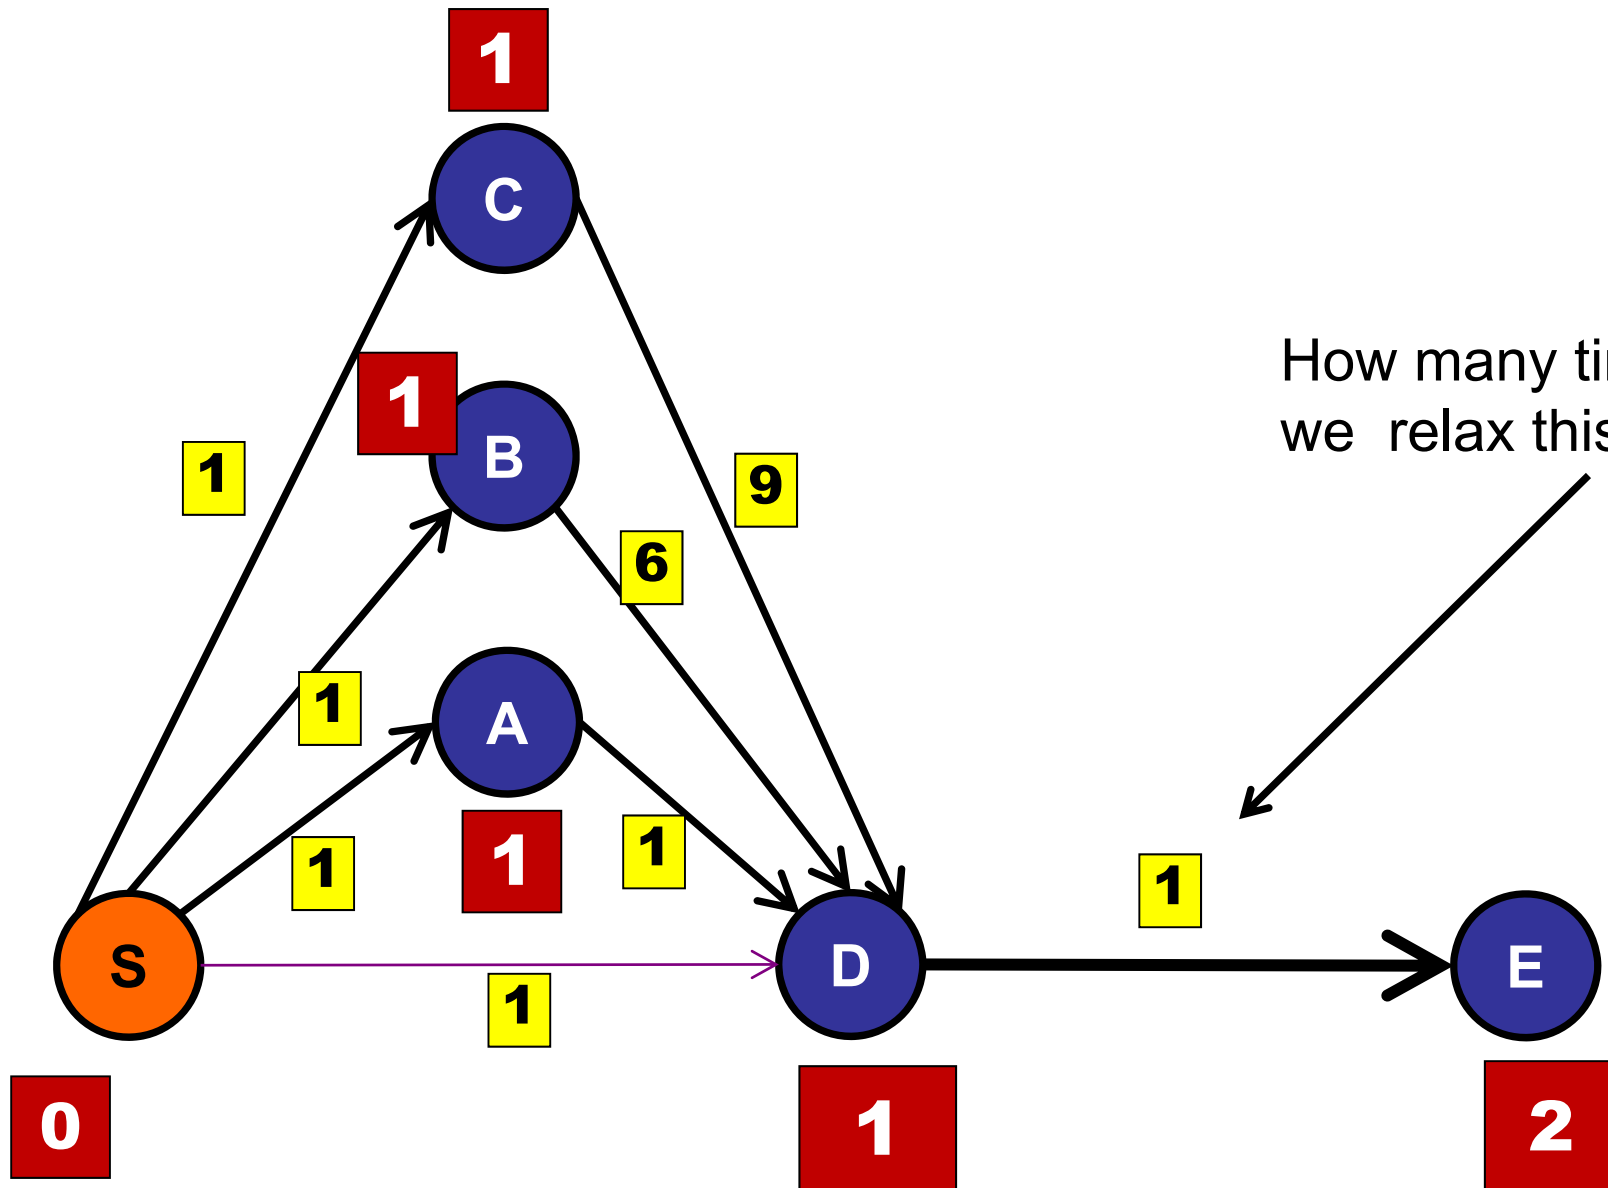


How many times might we relax this edge?

Shortest Paths




Shortest Paths



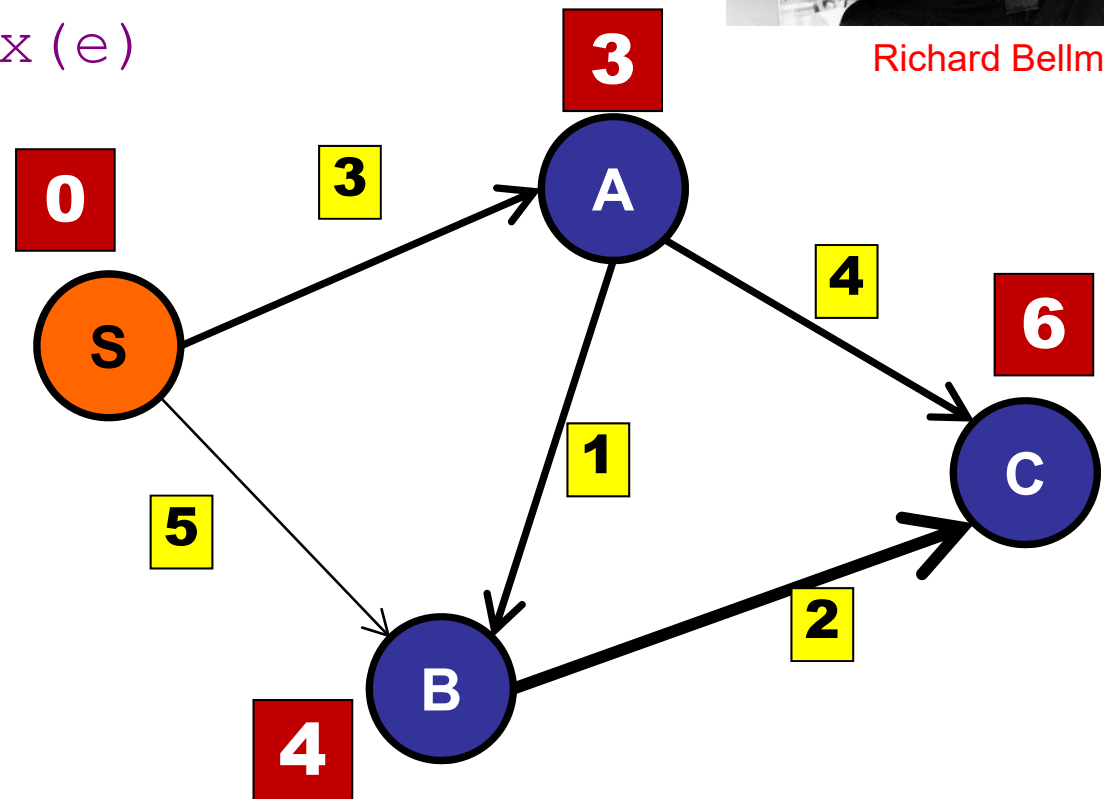
How many times might we relax this edge?

Bellman-Ford


```
n = V.length;  
for (i=0; i<n; i++)  
    for (Edge e : graph)  
        relax(e)
```



Richard Bellman




When can you terminate early?


1. When a relax operation has no effect.
2. When two consecutive relax operations have no effect.
-  3. When an entire sequence of $|E|$ relax operations have no effect.
4. Never. Only after $|V|$ complete iterations.

Bellman-Ford

```
n = V.length;  
for (i=0; i<n; i++)  
    for (Edge e : graph)  
        relax(e)
```



Richard Bellman



What is the running time of Bellman-Ford?

1. $O(V)$
2. $O(E)$
3. $O(V+E)$
4. $O(E \log V)$
- ✓ 5. $O(EV)$


Bellman-Ford

```
n = V.length;
```

```
for (i=0; i<n; i++)
```

```
    for (Edge e : graph)
```

```
        relax(e)
```




Bellman-Ford

Why does this work?


Bellman-Ford

Why does this work?



Bellman-Ford


Why does this work?



Look at minimum weight path from S to D.
(Path is simple: no loops.)

Bellman-Ford


Why does this work?



After 1 iteration, 1 hop estimate is correct.

Bellman-Ford

Why does this work?




After 1 iteration, 1 hop estimate is correct.

Shortest Paths


Key property:

If P is the shortest path from S to D , and if P goes through X , then P is also the shortest path from S to X (and from X to D).



Bellman-Ford


Why does this work?



After 1 iteration, 1 hop estimate is correct.

Bellman-Ford


Why does this work?



After 3 iterations, 3 hop estimate is correct.

Bellman-Ford


Why does this work?



After 4 iterations, D estimate is correct.


Bellman-Ford

What if edges have negative weight?



Bellman-Ford


What if edges have negative weight?



No problem!


Bellman-Ford

What if edges have negative weight?



Bellman-Ford


What if edges have negative weight?



$d(S,C)$ is infinitely negative!


Negative weight cycles

How to detect negative weight cycles?



Negative weight cycles

How to detect negative weight cycles?




Run Bellman-Ford for $|V|+1$ iterations.

If an estimate changes in the last iteration...
then negative weight cycle.


Bellman-Ford

Special case: all edges have the same weight



Bellman-Ford

Special case: all edges have the same weight.



Use regular Breadth-First Search.

Bellman-Ford Summary

Basic idea:

- Repeat $|V|$ times: relax every edge
- Stop when “converges”.
- $O(VE)$ time.

Special issues:

- If negative weight-cycle: impossible.
- Use Bellman-Ford to detect negative weight cycle.
- If all weights are the same, use BFS.

Faster algorithms?

Key idea:

Relax the edges in the “right” order.

Only relax each edge once:

- $O(E)$ cost (for relaxation step).

Necessary assumption:

All edges weights ≥ 0 .

Extending a path does not make it shorter!

Edsger W. Dijkstra


“Computer science is no more about computers than astronomy is about telescopes.”

“The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.”

“There should be no such thing as boring mathematics.”

“Elegance is not a dispensable luxury but a factor that decides between success and failure.”

“Simplicity is prerequisite for reliability.”




Edsger W. Dijkstra

“It is practically impossible to teach good programming to students that have had a prior exposure to BASIC: as potential programmers they are mentally mutilated beyond hope of regeneration.”

“The use of COBOL cripples the mind; its teaching should, therefore, be regarded as a criminal offense.”

“APL is a mistake, carried through to perfection. It is the language of the future for the programming techniques of the past: it creates a new generation of coding bums.”

“Object-oriented programming is an exceptionally bad idea which could only have originated in California.”



Faster algorithms?

Key idea:

Relax the edges in the “right” order.

Only relax each edge once:

- $O(E)$ cost (for relaxation step).


Necessary assumption:

All edges weights ≥ 0 .

Extending a path does not make it shorter!


Dijkstra's Algorithm (First Try)

Relax shortest edge first




Dijkstra's Algorithm (First Try)

Relax shortest edge first




Dijkstra's Algorithm (First Try)

Relax shortest edge first




Dijkstra's Algorithm (First Try)

Relax shortest edge first



Dijkstra's Algorithm (Failed Try)

Oops....




Dijkstra's Algorithm

Basic idea:


- Maintain distance estimate for every node.
- Begin with empty shortest-path-tree.
- Repeat:
 - Consider vertex with minimum estimate.
 - Add vertex to shortest-path-tree.
 - Relax all outgoing edges.

Shortest Paths



Dijkstra's Algorithm

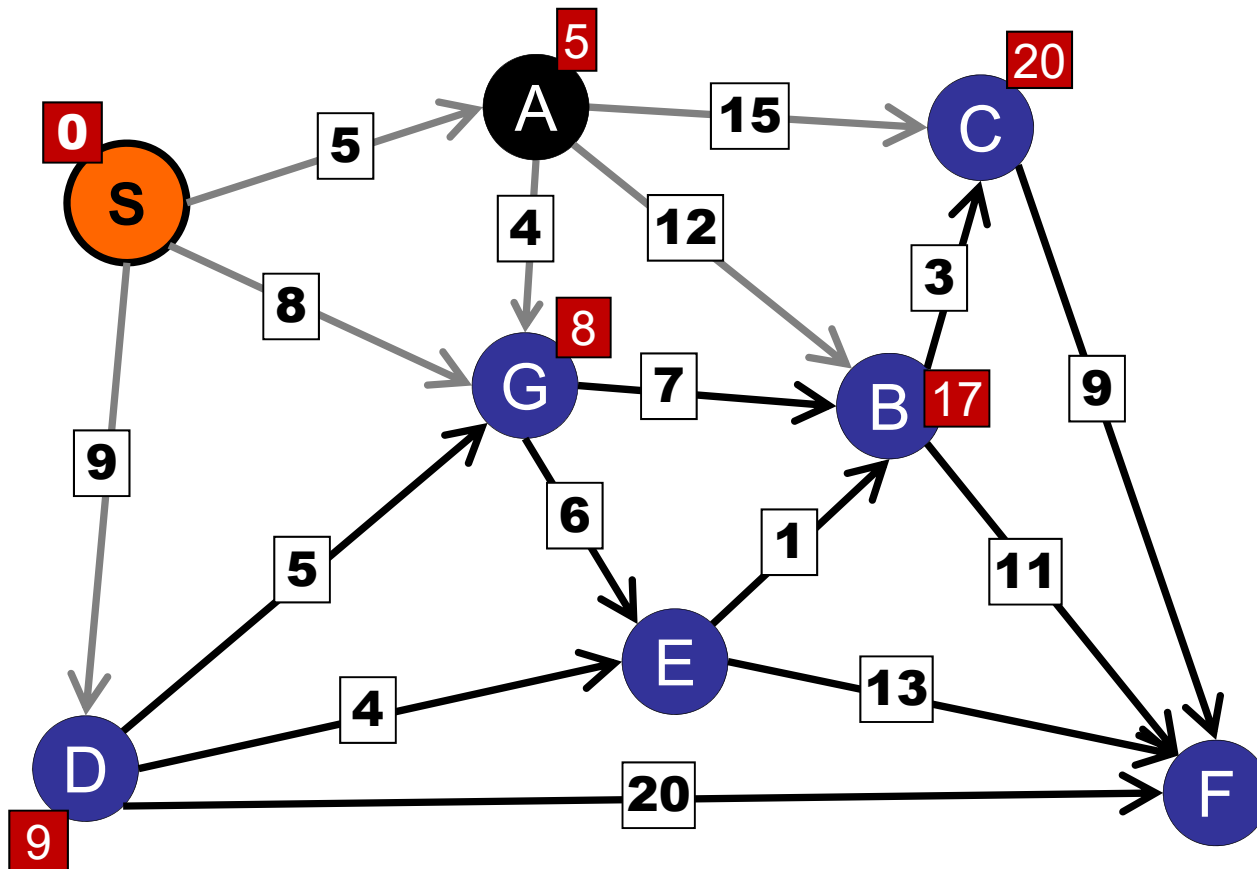
Step 2: Remove S and relax.



Vertex	Dist.
A	5
G	8
D	9

Dijkstra's Algorithm

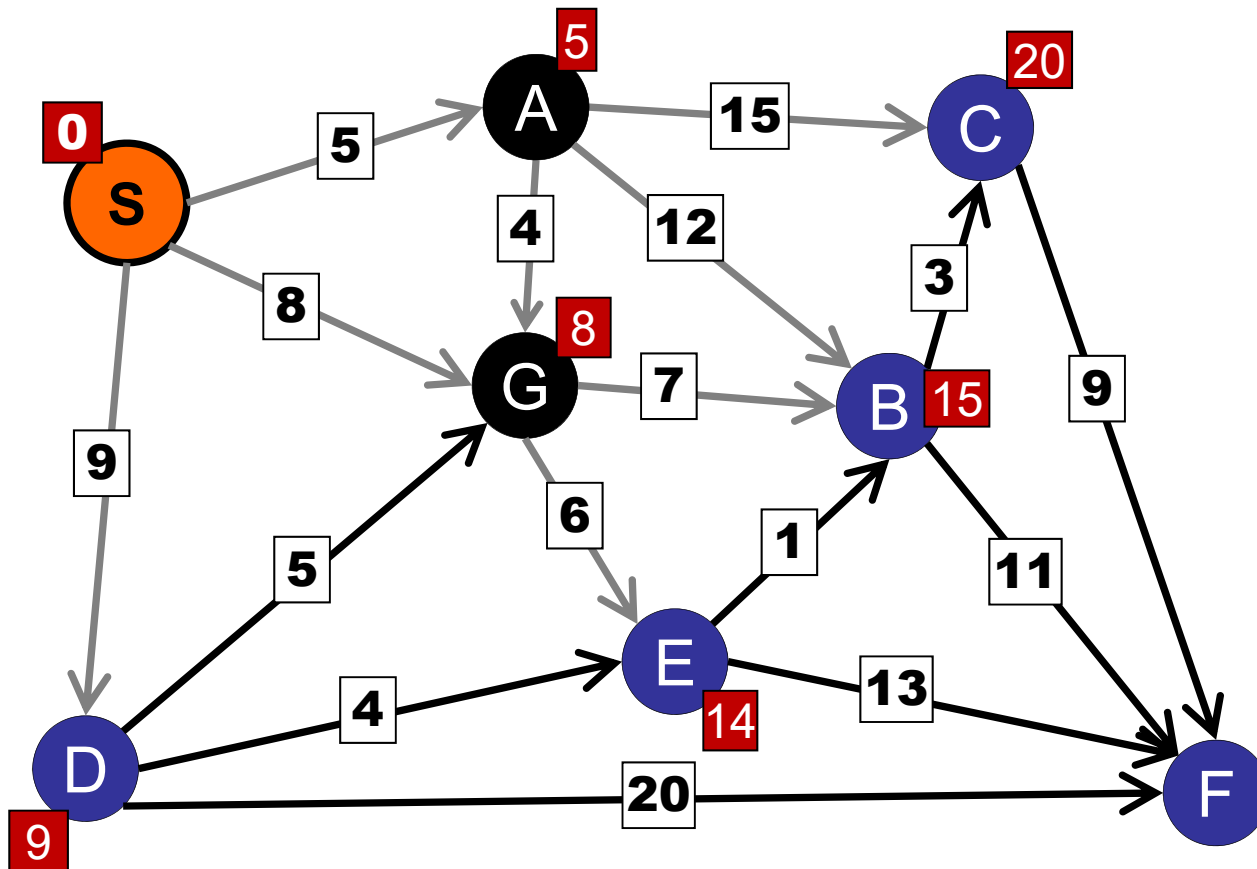
Step 3: Remove A and relax.



Vertex	Dist.
G	8
D	9
B	17
C	20

Dijkstra's Algorithm

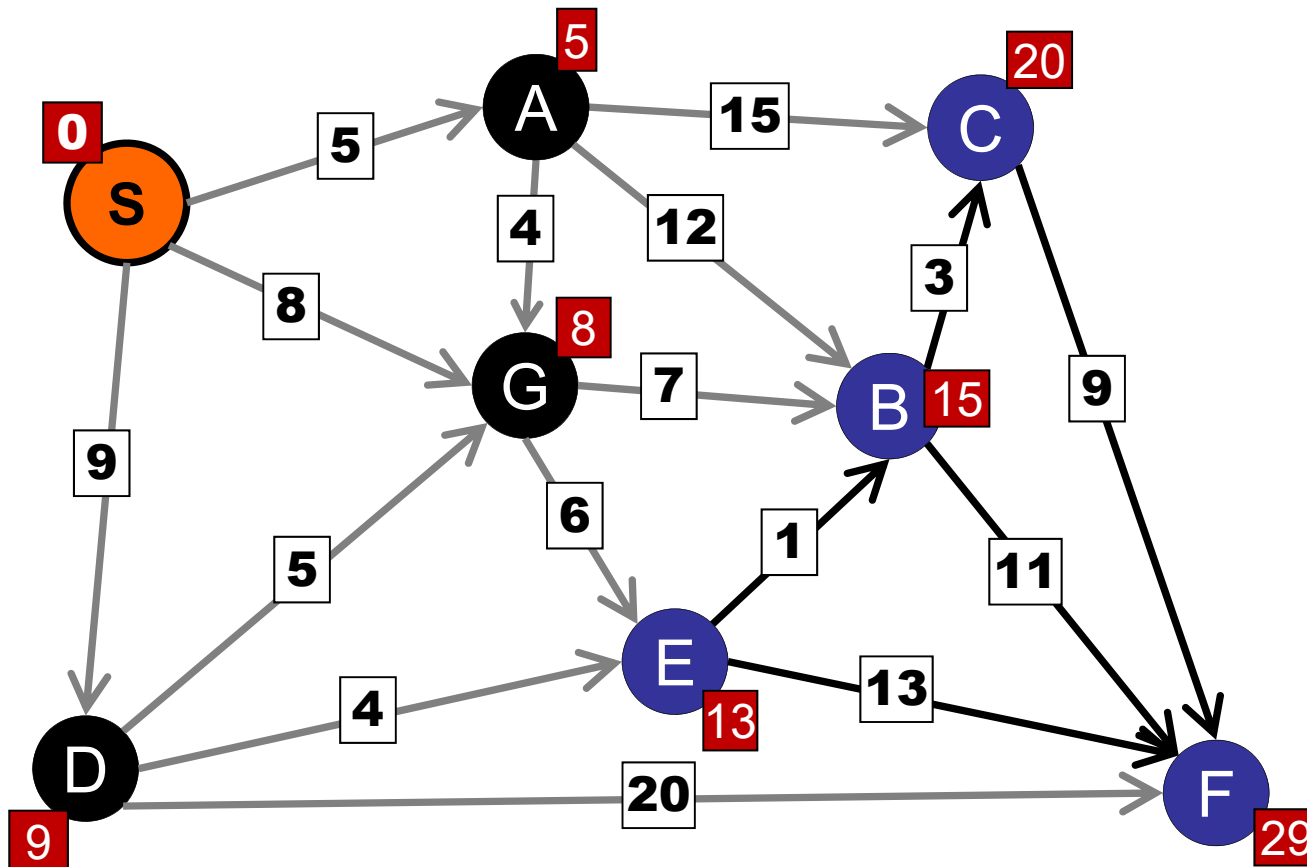
Step 4: Remove G and relax.



Vertex	Dist.
D	9
E	14
B	15
C	20

Dijkstra's Algorithm

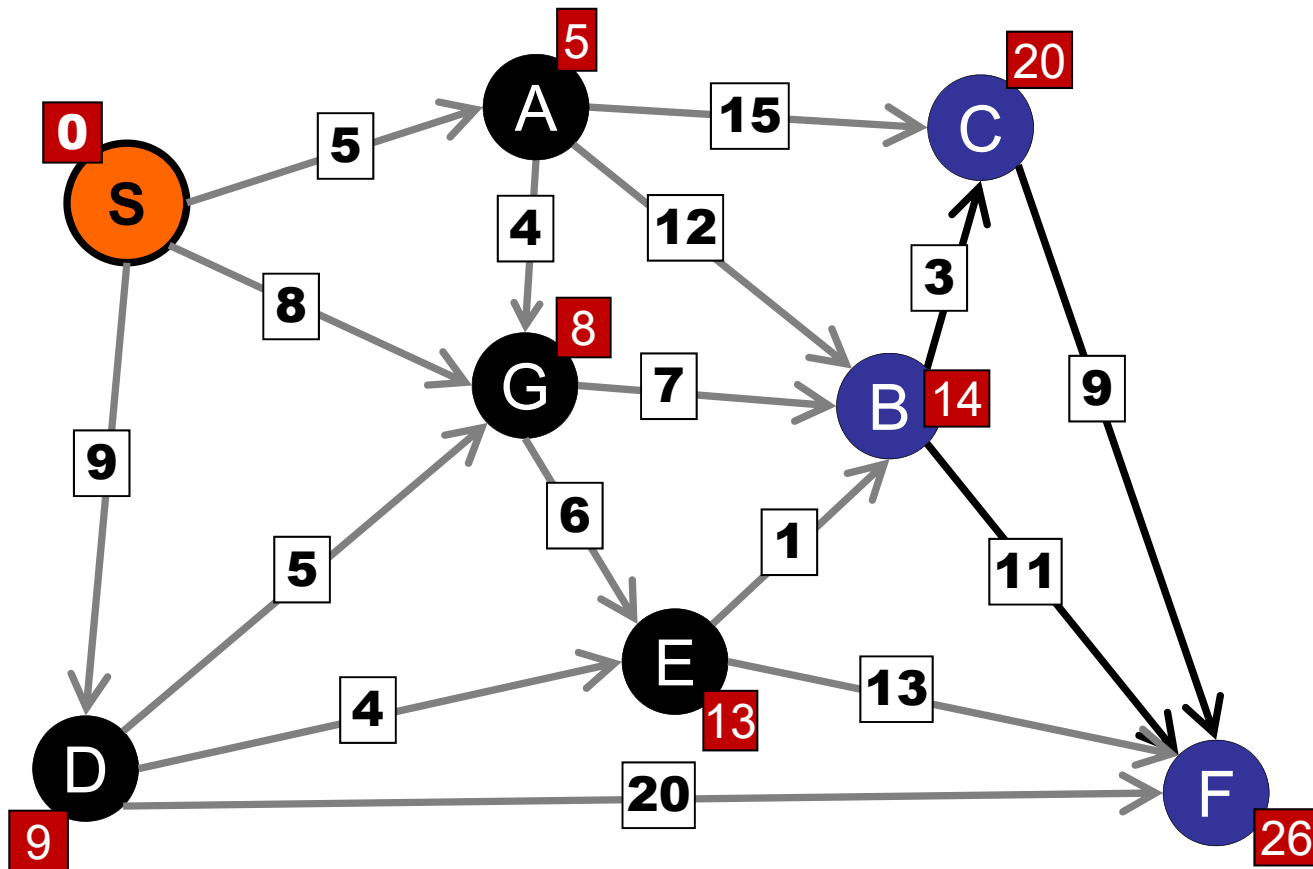
Step 5: Remove D and relax.



Vertex	Dist.
E	13
B	15
C	20
F	29

Dijkstra's Algorithm

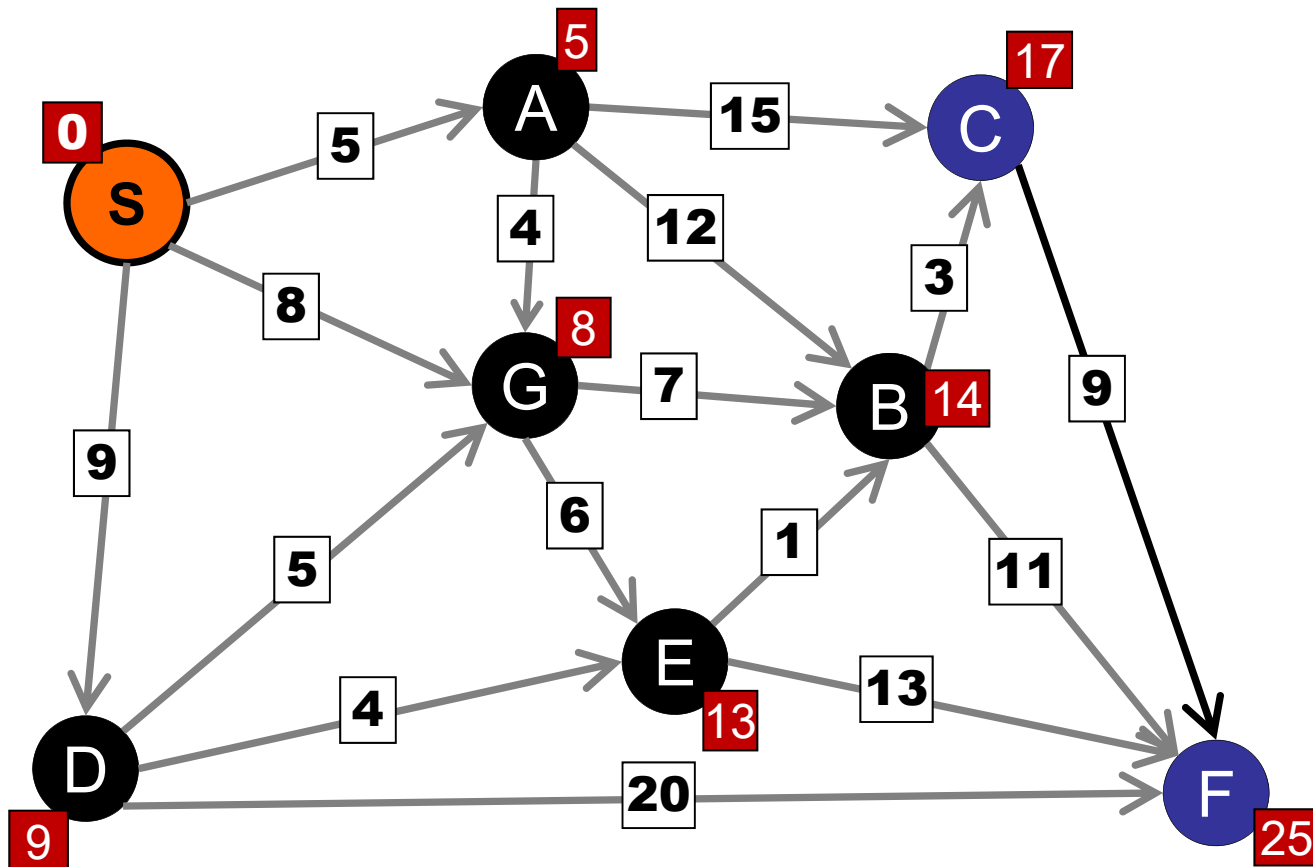
Step 5: Remove E and relax.



Vertex	Dist.
B	14
C	20
F	26

Dijkstra's Algorithm

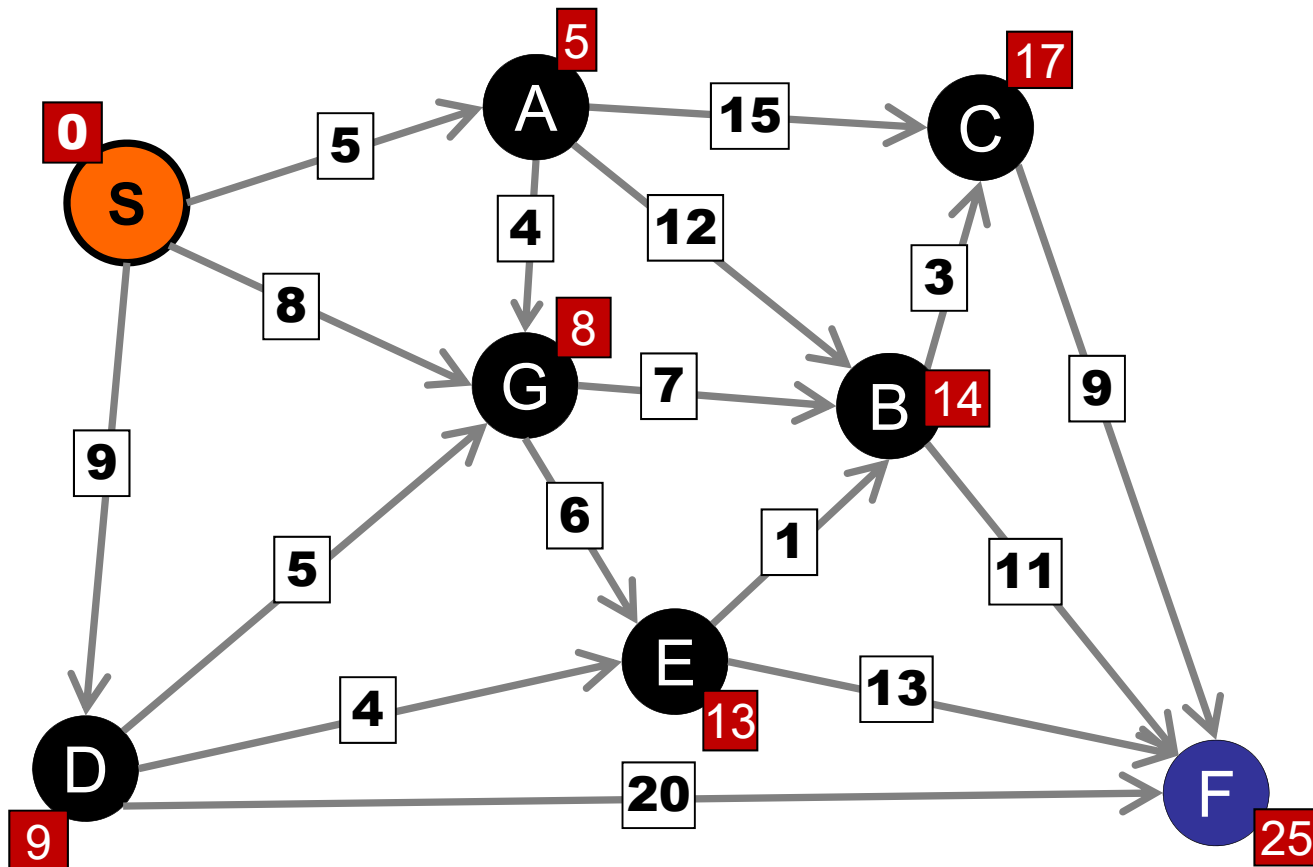
Step 5: Remove B and relax.



Vertex	Dist.
C	20
F	25

Dijkstra's Algorithm

Step 5: Remove C and relax.

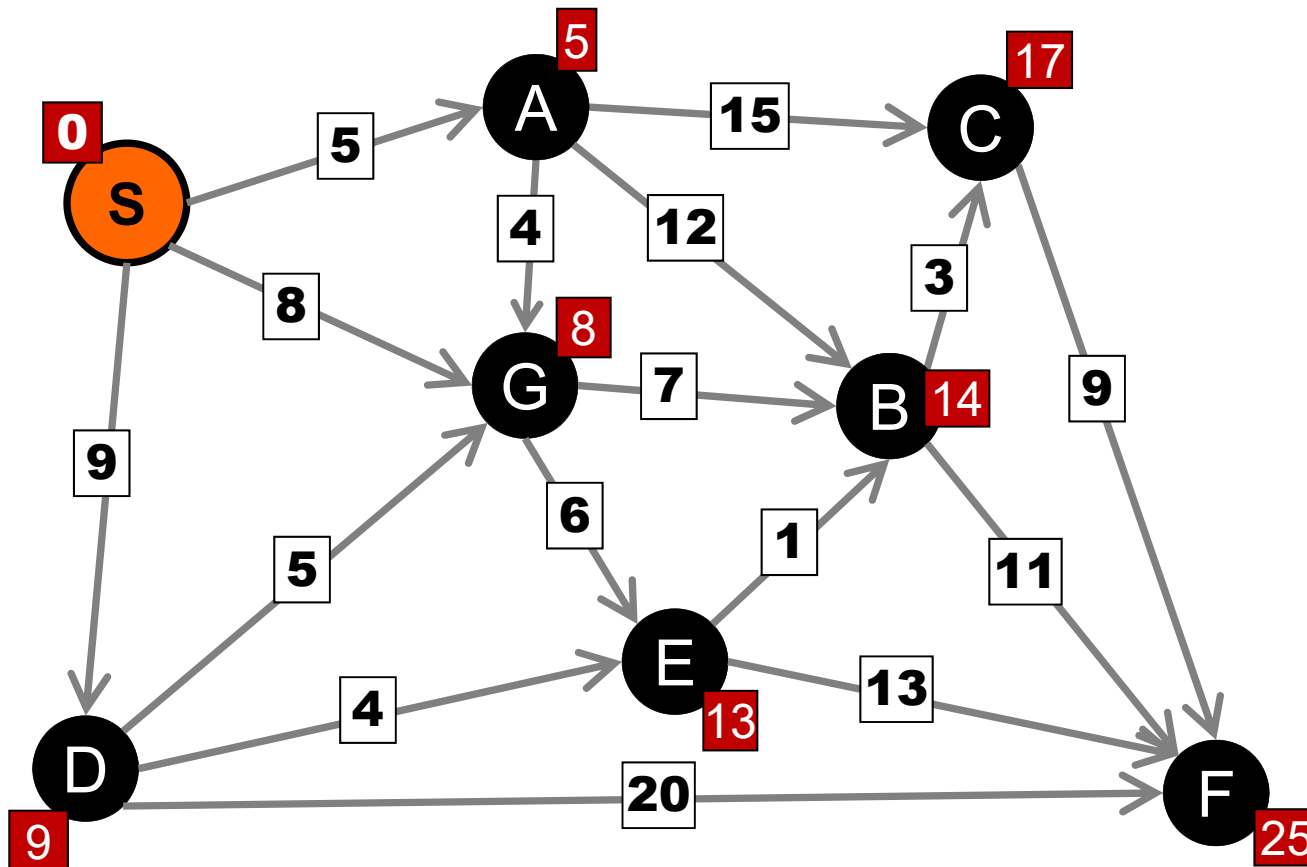


Vertex	Dist.
F	25

Dijkstra's Algorithm

Step 5: Remove F and relax.

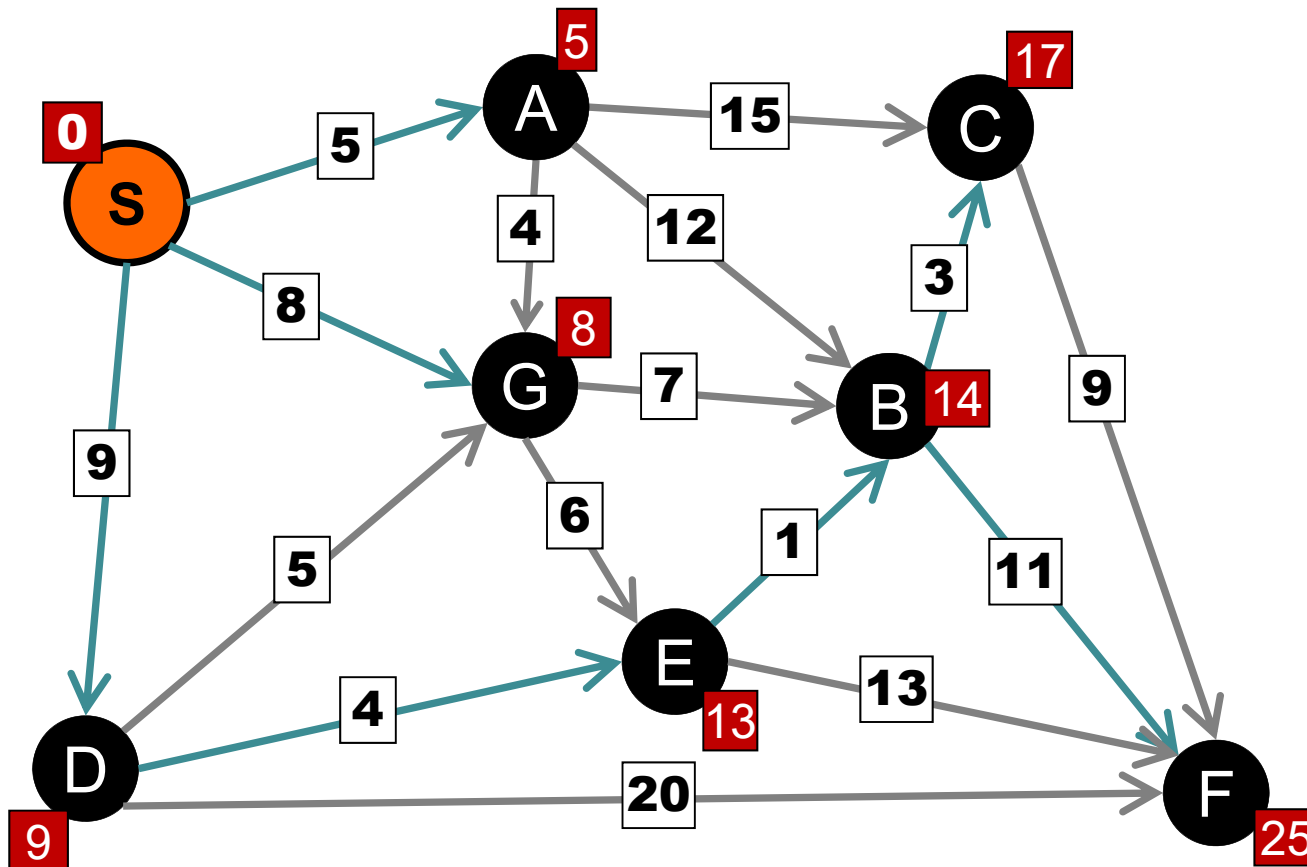
Vertex	Dist.



Dijkstra's Algorithm


Done

Vertex	Dist.



Dijkstra's Algorithm

Shortest Path Tree



Vertex	Dist.

What data structure to store vertices/distances?

1. Array
2. Linked list
3. Stack
4. Queue
- ✓ 5. AVL Tree
6. Huh?

Vertex	Dist.
B	14
C	20
F	26

Abstract Data Type

Priority Queue

interface **IPriorityQueue<Key, Priority>**

void insert(Key k, Priority p)

*insert k with
priority p*

Data extractMin()

*remove key with
minimum priority*

void decreaseKey(Key k, Priority p)

*reduce the priority of
key k to priority p*

boolean contains(Key k)

*does the priority
queue contain key k?*

boolean isEmpty()

*is the priority queue
empty?*

Notes:

Assume data items are unique.

```
public Dijkstra{
    private Graph G;
    private IPriorityQueue pq = new PriQueue();
    private double[] distTo;

    searchPath(int start) {
        pq.insert(start, 0.0);
        distTo = new double[G.size()];
        Arrays.fill(distTo, INFTY);
        distTo[start] = 0;
        while (!pq.isEmpty()) {
            int w = pq.deleteMin();
            for (Edge e : G[w].nbrList)
                relax(e);
        }
    }
}
```

Dijkstra's Algorithm

```
relax(Edge e) {  
    int v = e.from();  
    int w = e.to();  
    double weight = e.weight();  
    if (distTo[w] > distTo[v] + weight) {  
        distTo[w] = distTo[v] + weight;  
        parent[w] = v;  
        if (pq.contains(w))  
            pq.decreaseKey(w, distTo[w]);  
        else  
            pq.insert(w, distTo[w]);  
    }  
}
```

Abstract Data Type

Priority Queue

interface **IPriorityQueue<Key, Priority>**

void insert(Key k, Priority p)

*insert k with
priority p*

Data extractMin()

*remove key with
minimum priority*

void decreaseKey(Key k, Priority p)

*reduce the priority of
key k to priority p*

boolean contains(Key k)

*does the priority
queue contain key k?*

boolean isEmpty()

*is the priority queue
empty?*


Notes:

Assume data items are unique.

Priority Queue

AVL Tree


- Indexed by: priority
- Existing operations:
 - deleteMin()
 - insert(key, priority)



Priority Queue

AVL Tree


- Other operations:
 - contains()
 - decreaseKey(key, priority)
- How to find a vertex?



Priority Queue

AVL Tree

- Other operations:
 - contains()
 - decreaseKey(key, priority)
- Hash Table:
 - Map keys to locations in the binary tree.
 - Update hash table whenever the binary tree changes.



Dijkstra's Algorithm

Priority Queue by AVL tree:

- `insert(key, priority): $O(\log n)$`
- `deleteMin(): $O(\log n)$`
- `decreaseKey(key, priority): $O(\log n)$`
- `contains(key): $O(1)$`

What is the running time of Dijkstra's Algorithm, using an AVL tree Priority Queue?

1. $O(V + E)$
- ✓ 2. $O(E \log V)$
3. $O(V \log E)$
4. $O(V^2)$
5. $O(VE)$
6. None of the above

```
public Dijkstra{
    private Graph G;
    private MinPriQueue pq = new MinPriQueue();
    private double[] distTo;

    searchPath(int start) {
        pq.insert(start, 0.0);
        distTo = new double[G.size()];
        Arrays.fill(distTo, INFTY);
        distTo[start] = 0;
        while (!pq.isEmpty()) {
            int w = pq.deleteMin();
            for (Edge e : G[w].nbrList)
                relax(e);
        }
    }
}
```

How many times?

How many times?

Dijkstra's Algorithm

```
relax(Edge e) {
    int v = e.from();
    int w = e.to();
    double weight = e.weight();
    if (distTo[w] > distTo[v] + weight) {
        distTo[w] = distTo[v] + weight;
        parent[w] = v;
        if (pq.contains(w))
            pq.decreaseKey(w, distTo[w]);
        else
            pq.insert(w, distTo[w]);
    }
}
```

Dijkstra's Algorithm

Analysis:


- insert / deleteMin: $|V|$ times each
 - Each node is added to the priority queue **once**.
- relax / decreaseKey: $|E|$ times
 - Each edge is relaxed once.
- Priority queue operations: $O(\log V)$
- Total: $O((V+E)\log V) = O(E \log V)$

Dijkstra's Algorithm

Why does it work?

Dijkstra's Algorithm

Every edge crossing the boundary has been relaxed.




finished vertices:
distance is accurate.
Initially: just the source.

Dijkstra's Algorithm

Every edge crossing the boundary has been relaxed.

fringe vertices:
neighbor of a
finished vertex.

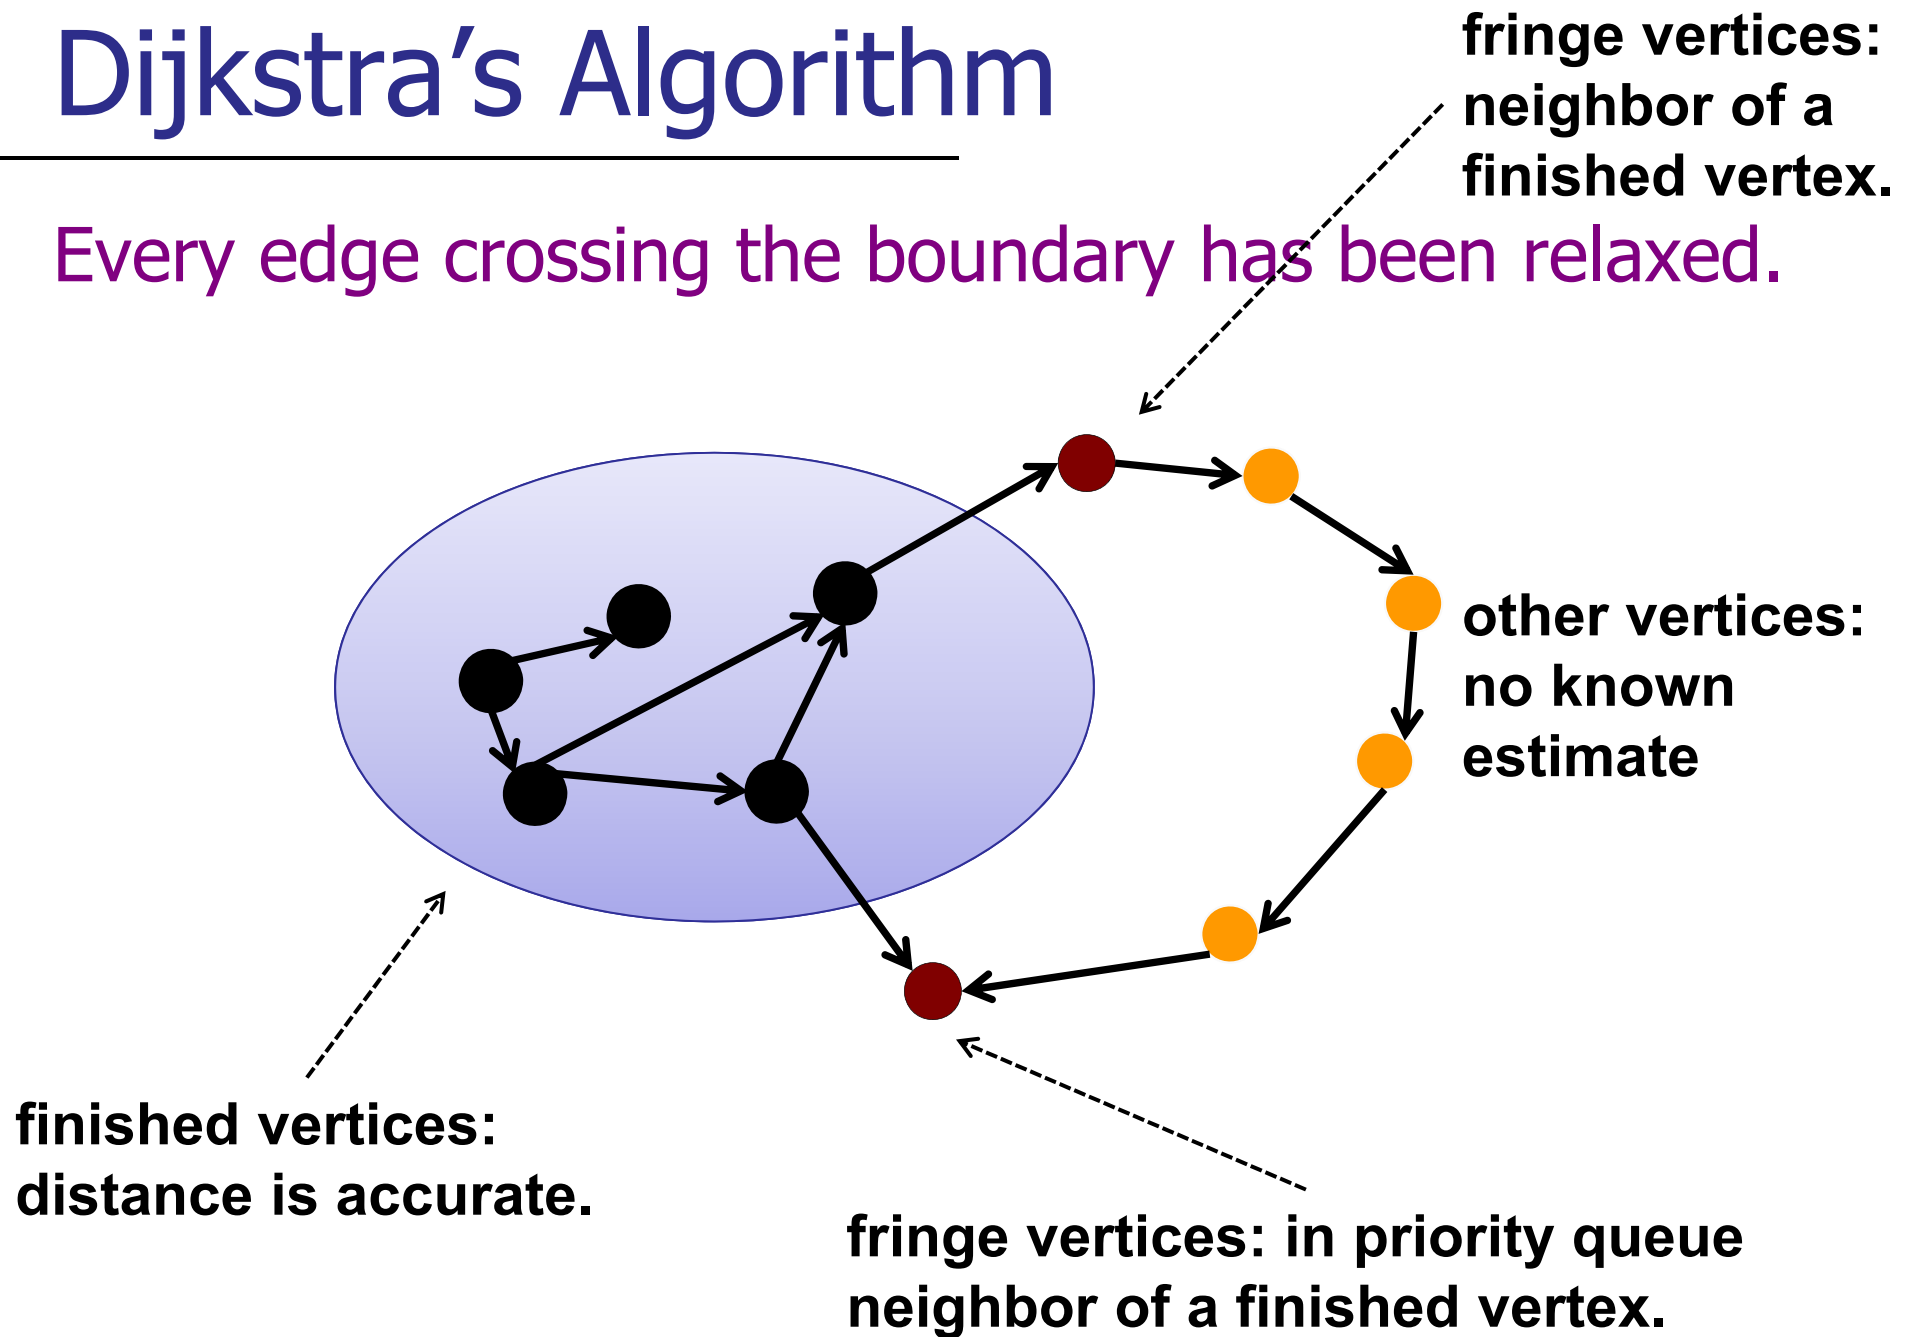


finished vertices:
distance is accurate.

fringe vertices: in priority queue
neighbor of a finished vertex.

Dijkstra's Algorithm

Every edge crossing the boundary has been relaxed.



Dijkstra's Algorithm

Proof by induction:

- Every “finished” vertex has correct estimate.
- Initially: only “finished” vertex is start.


Dijkstra's Algorithm

Proof by induction:

- Every “finished” vertex has correct estimate.
- Initially: only “finished” vertex is start.
- Inductive step:
 - Remove vertex from priority queue.
 - Relax its edges.
 - Add it to finished.
 - **Claim: it has a correct estimate.**


Dijkstra's Algorithm

Assume not: *fringe vertex is removed but not done*



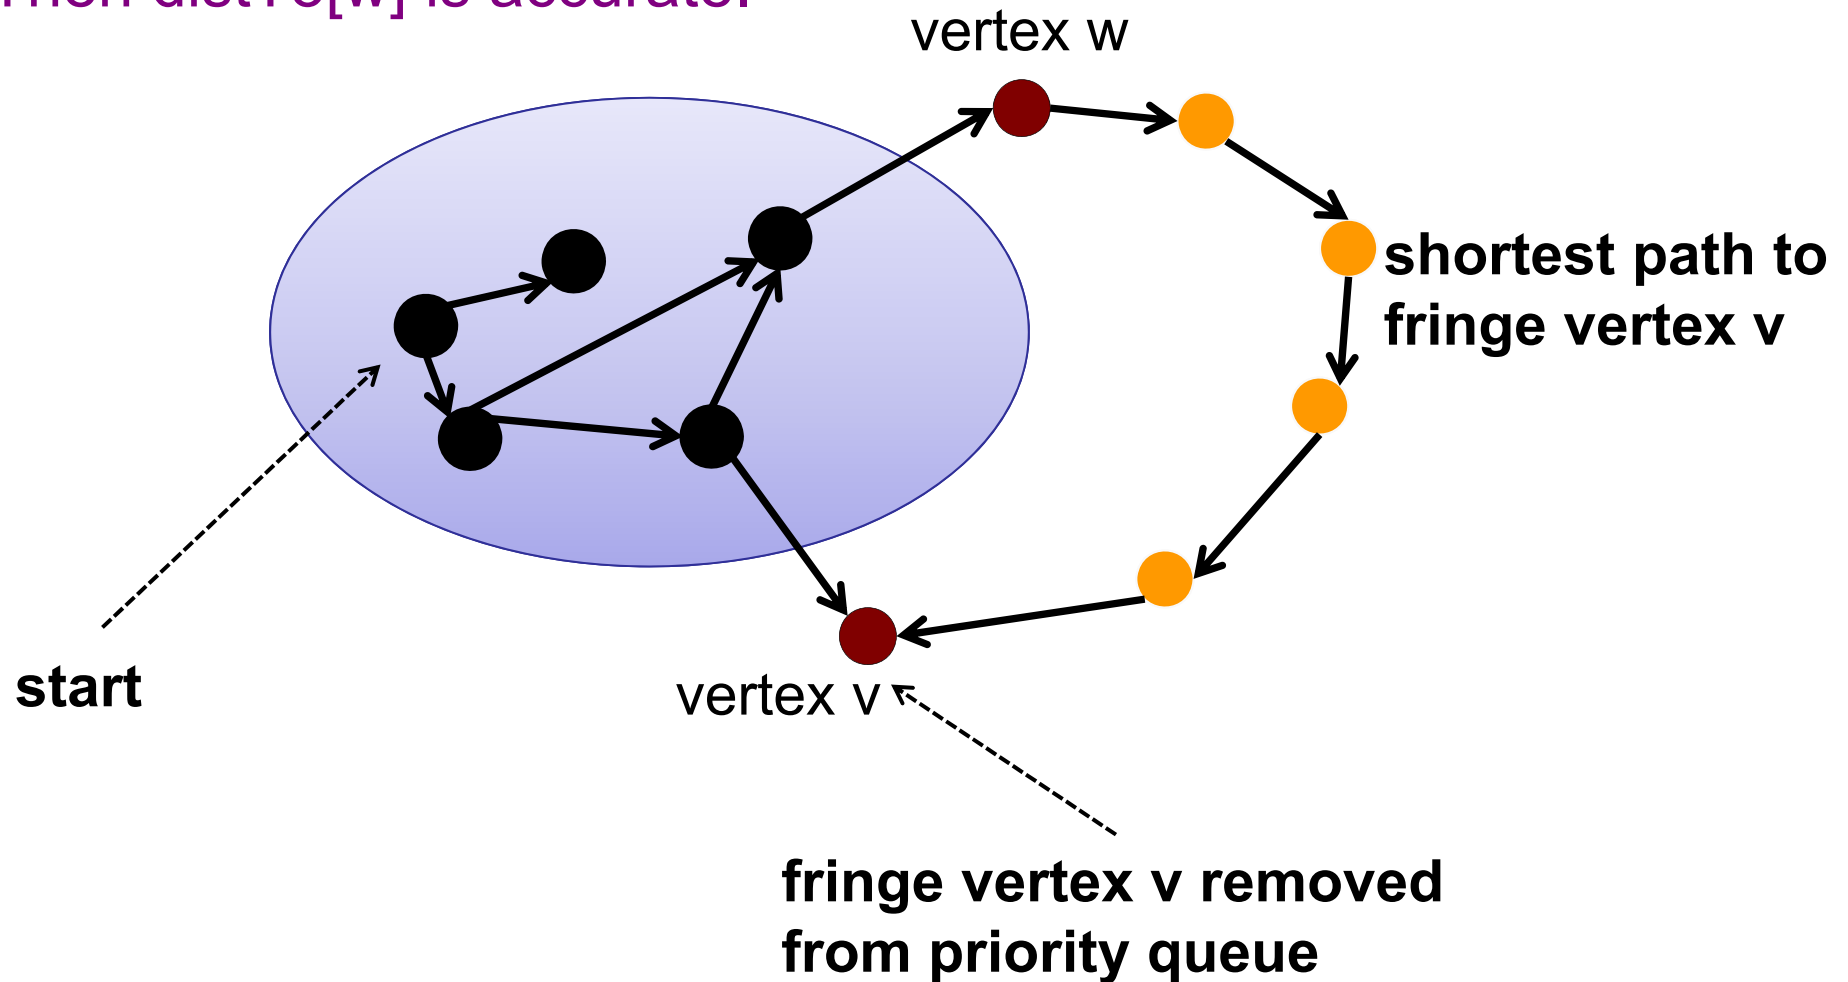
Dijkstra's Algorithm

Assume not: *fringe vertex has shorter path* **shortest path to fringe vertex**



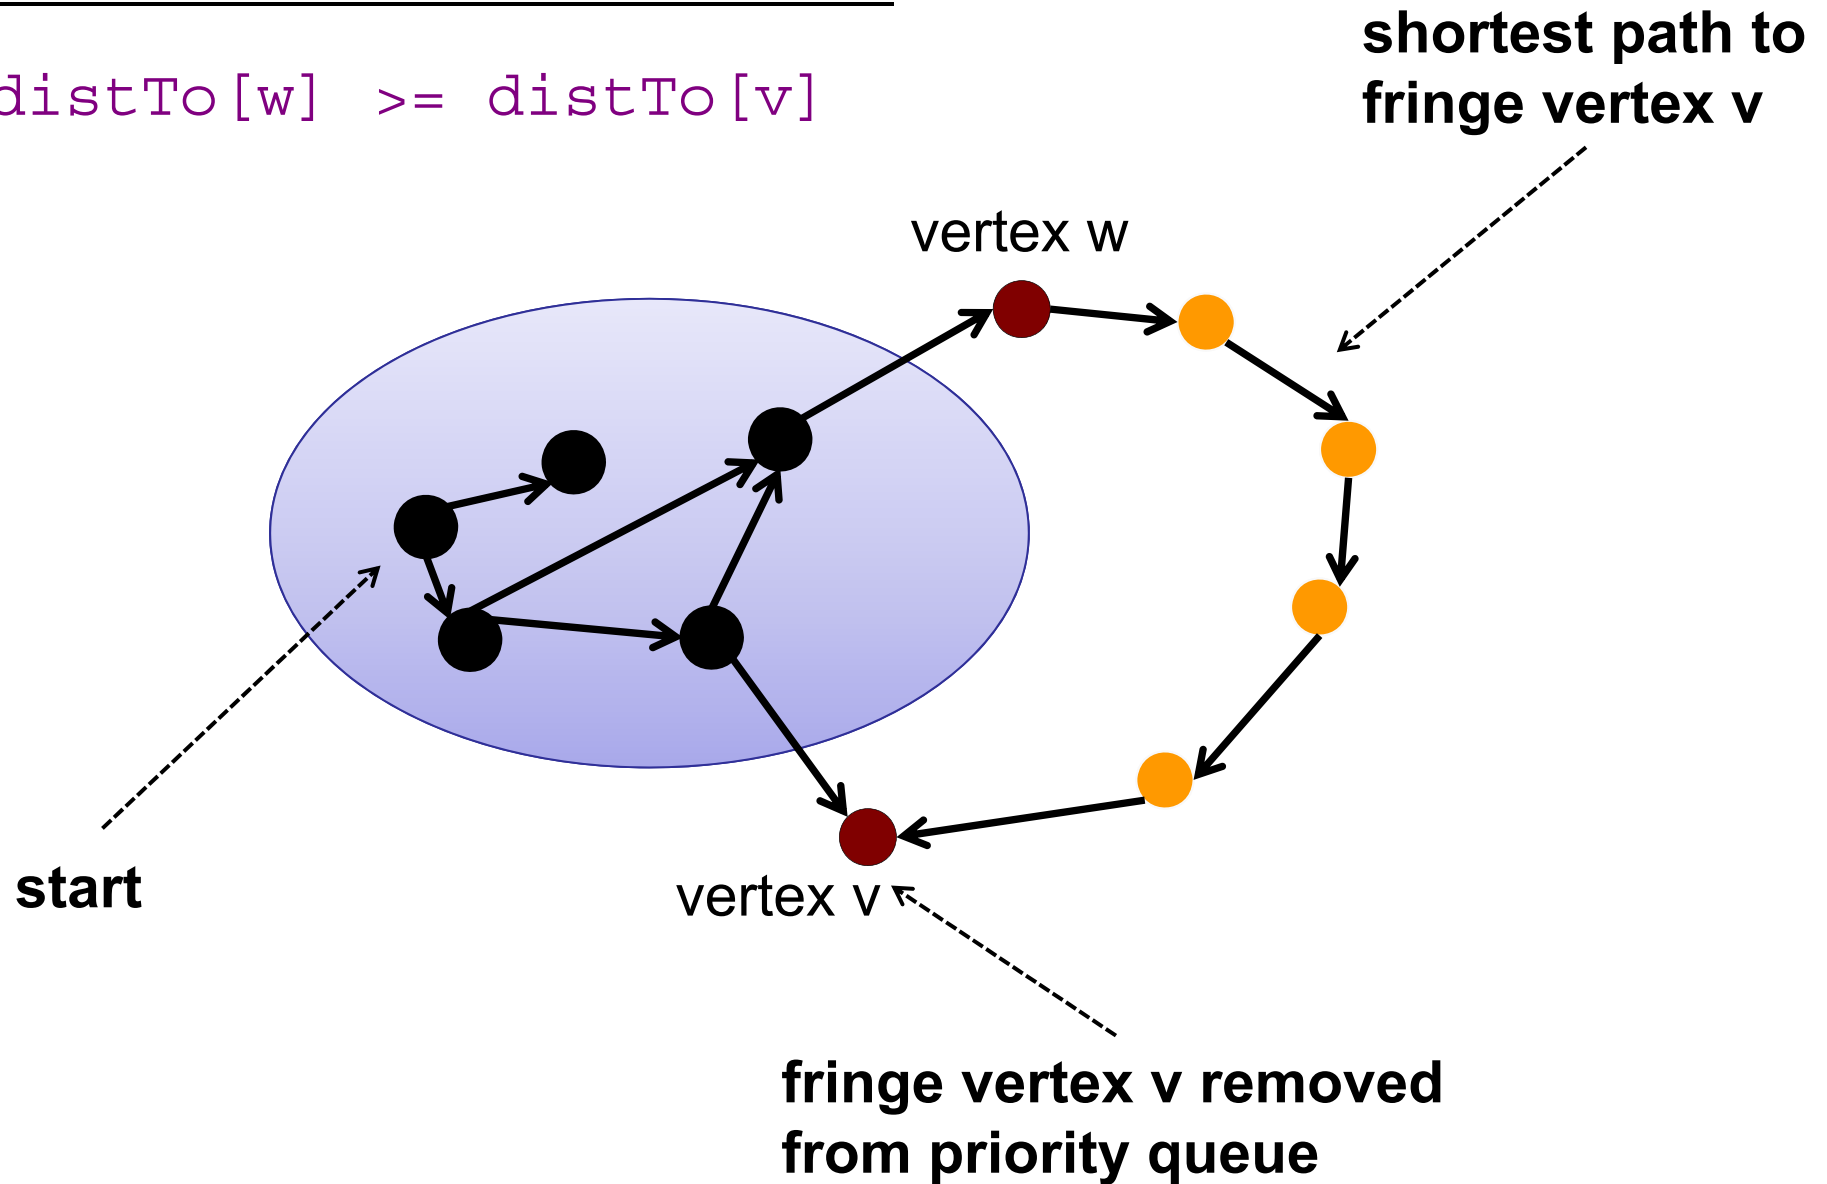
Dijkstra's Algorithm

If P is shortest path to v , then prefix of P is shortest path to w .
Then $\text{distTo}[w]$ is accurate.



Dijkstra's Algorithm

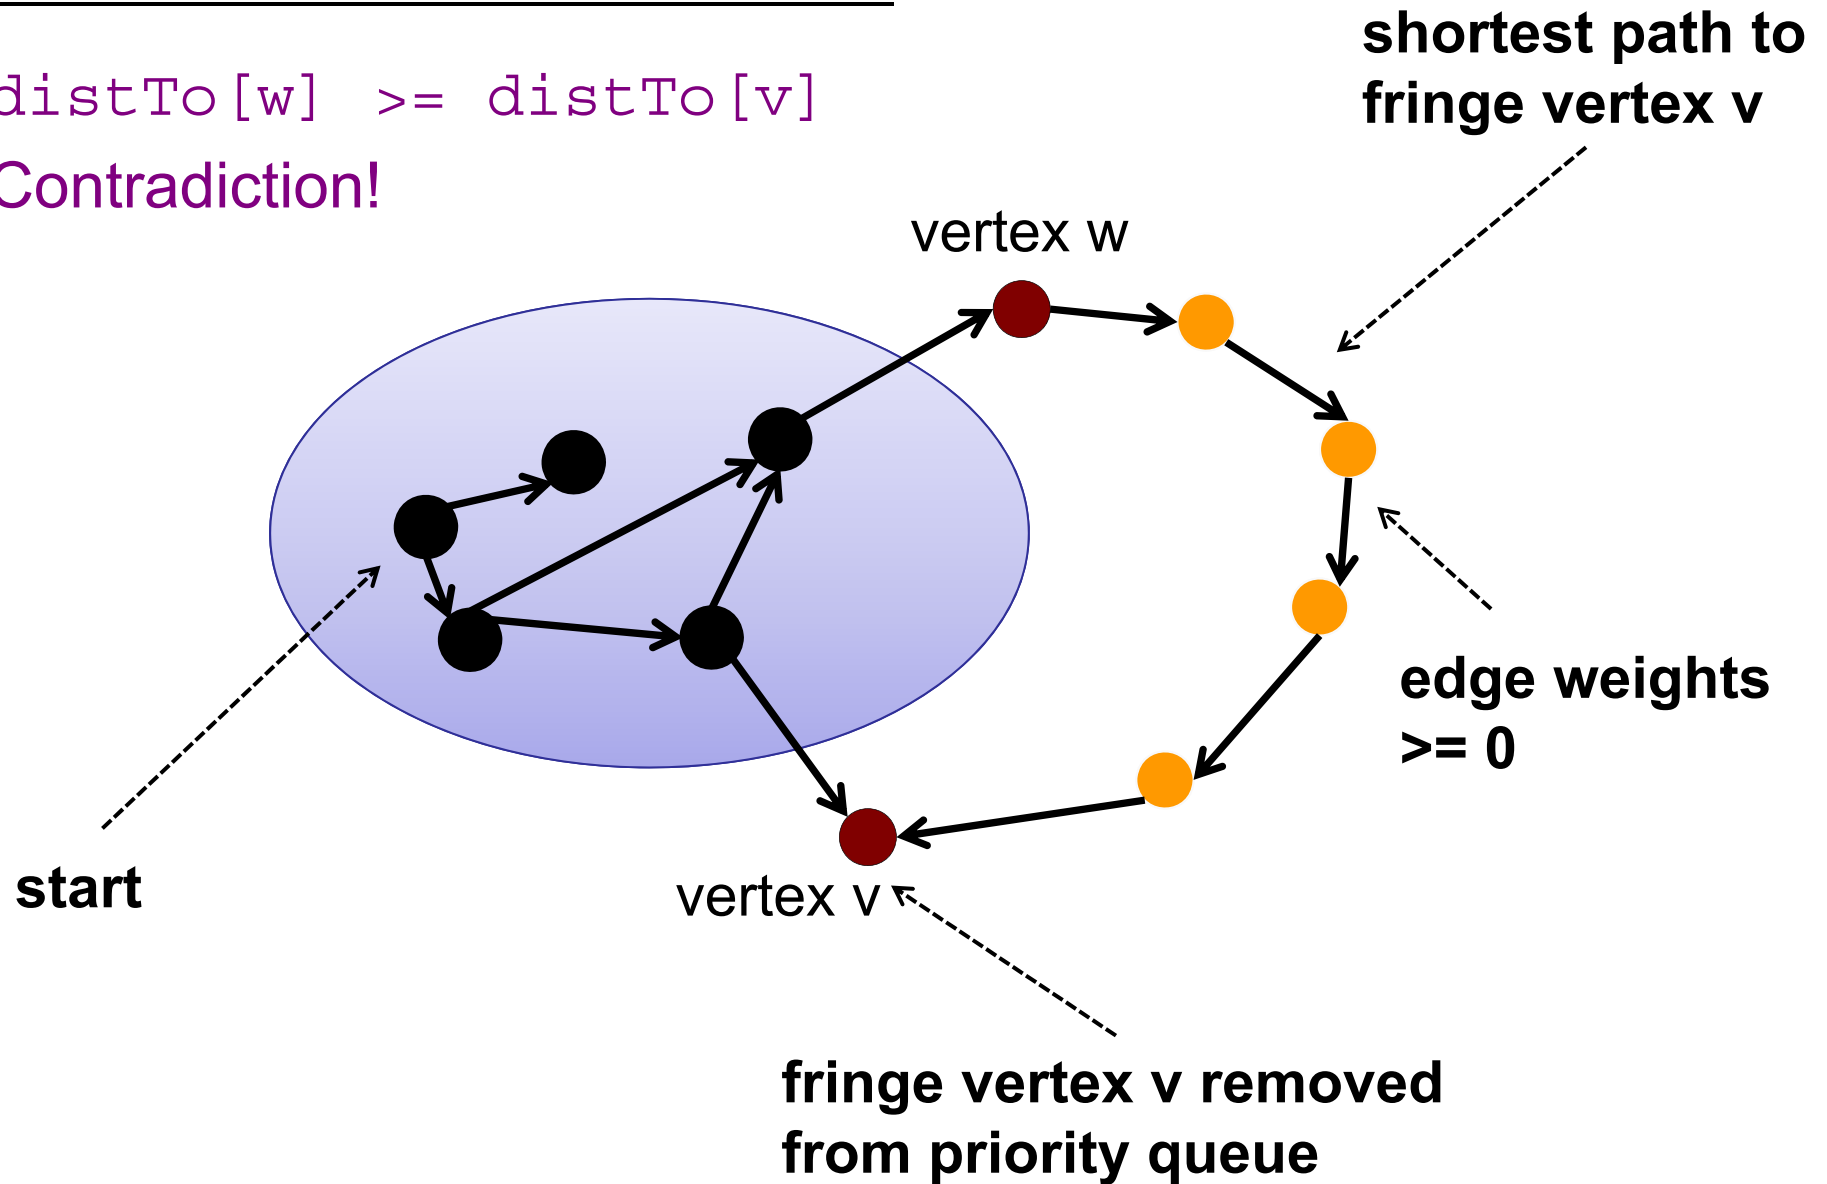
$\text{distTo}[w] \geq \text{distTo}[v]$



Dijkstra's Algorithm

$\text{distTo}[w] \geq \text{distTo}[v]$

Contradiction!



Dijkstra's Algorithm

Proof by induction:

- Every “finished” vertex has correct estimate.
- Initially: only “finished” vertex is start.
- Inductive step:
 - Remove vertex from priority queue.
 - Relax its edges.
 - Add it to finished.
 - **Claim: it has a correct estimate.**

Dijkstra's Algorithm

```
relax(Edge e) {  
    int v = e.from();  
    int w = e.to();  
    double weight = e.weight();  
    if (distTo[w] > distTo[v] + weight) {  
        distTo[w] = distTo[v] + weight;  
        parent[w] = v;  
        if (pq.contains(w))  
            pq.decreaseKey(w, distTo[w]);  
        else  
            pq.insert(w, distTo[w]);  
    }  
}
```

Extending a path does not make it shorter!

Dijkstra's Algorithm

Analysis:

- insert / deleteMin: $|V|$ times each
 - Each node is added to the priority queue **once**.
- decreaseKey: $|E|$ times
 - Each edge is relaxed once.
- Priority queue operations: $O(\log V)$
- Total: $O((V+E)\log V) = O(E \log V)$

Source-to-Destination Dijkstra

Can we stop as soon as we dequeue the destination?

- ✓ 1. Yes.
- 2. Only if the graph is sparse.
- 3. No.

Dijkstra's Algorithm

Source-to-Destination:

- What if you stop the first time you dequeue the destination?
- Recall:
 - a vertex is “finished” when it is dequeued
 - if the destination is finished, then stop

Dijkstra Summary

Basic idea:

- Maintain distance estimates.
- Repeat:
 - Find unfinished vertex with smallest estimate.
 - Relax all outgoing edges.
 - Mark vertex finished.
- $O(E \log V)$ time (with AVL tree).

Dijkstra's Performance


PQ Implementation	insert	deleteMin	decreaseKey	Total
Array	1	V	1	$O(V^2)$
AVL Tree	$\log V$	$\log V$	$\log V$	$O(E \log V)$
d-way Heap	$d \log_d V$	$d \log_d V$	$\log_d V$	$O(E \log_{E/V} V)$
Fibonacci Heap	1	$\log V$	1	$O(E + V \log V)$

Dijkstra Summary

Edges with negative weights?


Dijkstra's Algorithm

What goes wrong with negative weights?




Dijkstra's Algorithm

Edges with negative weights?



Dijkstra's Algorithm


Edges with negative weights?



Step 1: Remove A.
Relax A.
Mark A done.

Dijkstra's Algorithm

Edges with negative weights?



Step 1: Remove A.
Relax A.
Mark A done.


...

Step 4: Remove B.
Relax B.
Mark B done.

Oops: We need to
update A.

Dijkstra's Algorithm


What goes wrong with negative weights?



Dijkstra's Algorithm

Can we reweight?

e.g.: $\text{weight} += 10$




Can we reweight the graph?

1. Yes.
2. Only if there are no negative weight cycles.
- ✓ 3. No.

Dijkstra's Algorithm

Can we reweight?




Path S-B-A: 1

Path S-A: 4

Dijkstra's Algorithm

Can we reweight?



Path S-B-A: 21

Path S-A: 14

Dijkstra Summary

Basic idea:

- Maintain distance estimates.
- Repeat:
 - Find unfinished vertex with smallest estimate.
 - Relax all outgoing edges.
 - Mark vertex finished.
- $O(E \log V)$ time (with AVL tree Priority Queue).
- No negative weight edges!

Dijkstra Comparison

Same algorithm:

- Maintain a set of explored vertices.
 - Add vertices to the explored set by following edges that go from a vertex in the explored set to a vertex outside the explored set.
-
- **BFS:** Take edge from vertex that was discovered **least** recently.
 - **DFS:** Take edge from vertex that was discovered **most** recently.
 - **Dijkstra's:** Take edge from vertex that is **closest** to source.

Dijkstra Comparison

Same algorithm:

- Maintain a set of explored vertices.
 - Add vertices to the explored set by following edges that go from a vertex in the explored set to a vertex outside the explored set.
-
- BFS: Use queue.
 - DFS: Use stack.
 - Dijkstra's: Use priority queue.

Today

Single-Source Shortest Paths

- Weighted, Directed Graphs
- Bellman-Ford: simple, general
- Dijkstra: faster, only non-negative weights